

▶ Testing

[] ↳ 3 cells hidden

▼ ML Exercise 1-1

Question 1

▼ Loading Data

```
1 # Download dataset
2 !wget https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv

--2021-11-20 19:14:28-- https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, ..
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 12074 (12K) [text/plain]
Saving to: 'dataset1.csv.2'

dataset1.csv.2      100%[=====>]  11.79K  --.-KB/s    in 0s

2021-11-20 19:14:28 (86.2 MB/s) - 'dataset1.csv.2' saved [12074/12074]
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from tqdm.notebook import tqdm
5 import math
6
7 # Load the data using pandas
8 df = pd.read_csv("dataset1.csv")
```

```
1 # Show a sample of the data
2 df.head()
```

| | x | y |
|---|-----------|-----------|
| 0 | 0.097627 | 0.626964 |
| 1 | 0.430379 | 0.846452 |
| 2 | 0.205527 | 0.756017 |
| 3 | 0.089766 | 0.427504 |
| 4 | -0.152690 | -1.335228 |

```
1 # Show a description of the data (might be useful later)
2 df.describe()
```

| | x | y |
|-------|------------|------------|
| count | 300.000000 | 300.000000 |

▼ Helper functions

```
def hypothesis(X, theta):
```

▼ Loss Functions

```
def MSE(X, y, theta):
```

```
1 # h(theta) = theta transpose * X
2 def hypothesis(X, theta):
3     y1 = theta*X
4     return np.sum(y1, axis=1)
5
6 def MSE(X, y, theta):
7     y_hat = hypothesis(X, theta)
8     m = len(X)
9     return sum((y_hat-y)**2)/(2*m)
10
11 def RMSE(X, y, theta):
12     y_hat = hypothesis(X, theta)
13     m = len(X)
14     return np.sqrt(sum((y_hat-y)**2)/(2*m))
15
16 def MAE(X, y, theta):
17     y_hat = hypothesis(X, theta)
18     m = len(X)
19     return sum(np.abs((y_hat-y)))/(2*m)
20
21 # Loss functions Derivatives
22 def MSE_prim(X, y, i, theta):
23     y_hat = hypothesis(X, theta)
24     Xi = X.iloc[:, i]
25     m = len(X)
26     return sum((y_hat-y) * Xi) / m
27
28 def RMSE_prim(X, y, i, theta):
29     # src : https://math.stackexchange.com/questions/4065532/rmse-derivatives
30     mse = MSE(X, y, theta)
31     mse_prim = MSE_prim(X, y, i, theta)
32
33     return mse_prim / 2 / np.sqrt(mse)
34
35 def MAE_prim(X, y, i, theta):
36     # src : https://stats.stackexchange.com/questions/312737/mean-absolute-error-mae-derivative
37     # src2 : https://github.com/chenxingwei/machine_learning_from_scratch/blob/master/algorithm/2.linearRegressionGradientDescent.py
38     y_hat = hypothesis(X, theta)
39     # print(np.sum((X.T*(np.sign(y_hat-y)/len(X))), axis=1)[i])
40     return np.sum((X.T*(np.sign(y_hat-y)/len(X))), axis=1)[i]
41
```

▼ Gradient Descent

```
1 def gradientDescent(X, y, theta, lr, iteration, X_valid, y_valid, loss_fn = MSE, loss_fn_prim = MSE_prim, decay=0.0):
2     # Training loss per iteration history
3     train_loss_history = []
4     # Validation loss per iteration history
5     validation_loss_history = []
6     # weights progression towards the optimal value
7     theta_history = []
8
9     # Progress bar
10    with tqdm(total=iteration) as pbar:
11        for iter in range(iteration):
12            for i in range(0, len(X.columns)):
13                # partial derivative of loss function with respect to Xi
14                gradient = loss_fn_prim(X, y, i, theta)
15
16                # TODO: Learning rate decay
17                # lr = lr * 1/(1 + decay * iter)
```

```

18
19     # Actual "Gradient Descent" !
20     theta[i] -= lr * gradient
21
22     # Calculating the loss after each iteration
23     # of updating the weights using Gradient Descent
24     loss = loss_fn(X, y, theta)
25     if X_valid is not None and y_valid is not None:
26         validation_loss = loss_fn(X_valid, y_valid, theta)
27
28     # Save the history of loss and weights
29     train_loss_history.append(loss)
30     if X_valid is not None and y_valid is not None:
31         validation_loss_history.append(validation_loss)
32     theta_history.append(theta.copy())
33
34     # Update progress bar
35     pbar.update(1)
36
37     history = {"train_loss":train_loss_history,
38               "validation_loss":validation_loss_history,
39               "weights":theta_history}
40     # returns loss history, latest loss, weights
41     return history, loss, theta

```

```

1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6 def polynomial_to_linear_regression(X, polynomial_degree):
7     for i in range(2, 2 + polynomial_degree):
8         X['x'+str(i)] = X['x']**i

```

▼ Normal Equation

```

1 def normalEquation(X, y):
2     #  $(X^T X)^{-1} X^T Y$ 
3     XTX = np.dot(X.T,X)
4     XTX_inverse = np.linalg.inv(XTX)
5     XTY = np.dot(X.T,y)
6     theta = np.dot(XTX_inverse, XTY)
7     return theta
8
9 def regularizedNormalEquation(X, y, lambd=0.1):
10    #  $(X^T X + \lambda I)^{-1} X^T Y$ 
11    XTX = np.dot(X.T,X) + np.dot(np.identity(X.shape[1]),lambd)
12    XTX_inverse = np.linalg.inv(XTX)
13    XTY = np.dot(X.T,y)
14    theta = np.dot(XTX_inverse, XTY)
15    return theta

```

▼ Plotting related

```

1 # helper function used to plot a polynomial
2 def polyCoefficients(x, coeffs):
3     o = len(coeffs)
4     y = 0
5     for i in range(o):
6         y += coeffs[i]*x**i
7     return y

```

```

1 # Plots a polynomial on top of the original data
2 def plot_curve(X, y, theta, c='r', title='', resolution=100):
3     plt.figure()
4     plt.title(title)
5     # Plot the original data
6     plt.scatter(x=X['x'],y= y)

```

```

7
8 x = np.linspace(-1, 1, resolution)
9 # Plot the fitted polynomial over the data
10 plt.plot(x, polyCoefficients(x, theta), c=c, linewidth=4)
11
12 plt.show()


```

▼ Part 1 - Plotting the data

```

1 # Plot the data using matplotlib
2 df.plot(kind='scatter', x='x', y='y')
3 plt.show()

```



▼ Part 2 - Shuffle

```

1 def plot_colorize(df):
2     '''
3     Assigns 'red' color to the first half of the data
4     and 'blue' to the rest
5
6     If the data is well shuffled we should see random red
7     and blue circles everywhere.
8
9     If the data is NOT well shuffled we might see a pattern between
10    circles' position and their color.
11    '''
12    col = {True:"red", False:"blue"}
13    colors = list(map(lambda x:col[x<df.shape[0]/2],df.index.tolist()))
14
15    df.plot(kind='scatter', x='x', y='y', color=colors)
16    plt.show()

```

▼ Default data


```

1 plot_colorize(df)
2 print(" *9,"Data seems to be well shuffled.")

```

Sorted data


```
1 # Sort the data based on 'x' first, then do the
2 # previous part to see the result.
3 sorted_df = df.sort_values(by='x', ascending=True, ignore_index=True)
4
5 plot_colorize(sorted_df)
6 print(" *6,"Data doesn't seem to be well shuffled.")
```



Data doesn't seem to be well shuffled.

Shuffled data

```
1 # Let's shuffle the data anyways (just in case)
2
3 # pandas Doc: specifying drop=True prevents .reset_index()
4 # from creating a column containing the old index entries.
5 shuffled_df = df.sample(frac=1).reset_index(drop=True)
6
7 plot_colorize(shuffled_df)
8 print(" *9,"Data seems to be well shuffled.")
```



Data seems to be well shuffled.

Part 3 - Gradient Descent

Finding The optimal "Theta" values (weights)

Data Prepration

Add a column for bias

```
1 # Add a new column for simplicity of the calculations
2 # acts as the bias term
```

```
3 shuffled_df = pd.concat([pd.Series(1, index=shuffled_df.index, name='0'), shuffled_df], axis=1)
4 shuffled_df.head()
```

| | 0 | x | y |
|---|---|-----------|-----------|
| 0 | 1 | 0.963659 | -0.527811 |
| 1 | 1 | 0.312659 | 1.348749 |
| 2 | 1 | -0.125936 | -0.631524 |
| 3 | 1 | 0.341276 | -0.078979 |
| 4 | 1 | 0.923873 | 0.217419 |

▼ Seperate X,y

```
1 # Split training data into X and y
2 X = shuffled_df.drop(columns="y")
3 y = shuffled_df.iloc[:, 2]
4
5 print(X.head(),end="\n\n")
6 print(y.head())
```

| | 0 | x |
|---|---|-----------|
| 0 | 1 | 0.963659 |
| 1 | 1 | 0.312659 |
| 2 | 1 | -0.125936 |
| 3 | 1 | 0.341276 |
| 4 | 1 | 0.923873 |

| | |
|---|-----------|
| 0 | -0.527811 |
| 1 | 1.348749 |
| 2 | -0.631524 |
| 3 | -0.078979 |
| 4 | 0.217419 |

Name: y, dtype: float64

```
1 # Split to train and valid
2 split = 0.7
3
4 X_train = X.iloc[ : int(len(X)*split),:].reset_index(drop=True)
5 X_valid = X.iloc[int(len(X)*split) : ,:].reset_index(drop=True)
6
7 y_train = y.iloc[ : int(len(X)*split)].reset_index(drop=True)
8 y_valid = y.iloc[int(len(X)*split) : ].reset_index(drop=True)
9
10 print(f"Train X size = {len(X_train)}")
11 print(f"Train y size = {len(y_train)}")
12 print(f"Valid X size = {len(X_valid)}")
13 print(f"Valid y size = {len(y_valid)}")
```

```
Train X size = 210
Train y size = 210
Valid X size = 90
Valid y size = 90
```

```
1 # Save a copy of X and y
2 # TODO might not need it
3 X_train_org = X_train.copy()
4 y_train_org = y_train.copy()
```

▼ Polynomial Regression

a basic example

```
1 # polynomial degree
2 polynomial_degree = 12 #@param {type: "number"}
3 learning_rate = 0.030 #@param {type: "number"}
4 iterations = 1000 #@param {type: "number"}
5
```

polynomial_degree: 12

learning_rate: 0.030

iterations: 1000

Convert

$aX + bX^2 + cX^3 + d$

to

$aX1 + bX2 + cX3 + d$

```
1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6 polynomial_to_linear_regression(X_train, polynomial_degree)
7 polynomial_to_linear_regression(X_valid, polynomial_degree)
8
9 X_train.head()
```

| | 0 | x | x2 | x3 | x4 | x5 | x6 | x7 | x8 | x9 | x10 | x11 |
|---|---|-----------|----------|-----------|----------|-----------|----------|---------------|--------------|---------------|--------------|---------------|
| 0 | 1 | 0.963659 | 0.928638 | 0.894890 | 0.862369 | 0.831029 | 0.800829 | 7.717257e-01 | 7.436803e-01 | 7.166540e-01 | 6.906099e-01 | 6.655123e-01 |
| 1 | 1 | 0.312659 | 0.097756 | 0.030564 | 0.009556 | 0.002988 | 0.000934 | 2.920776e-04 | 9.132075e-05 | 2.855227e-05 | 8.927129e-06 | 2.791149e-06 |
| 2 | 1 | -0.125936 | 0.015860 | -0.001997 | 0.000252 | -0.000032 | 0.000004 | -5.024022e-07 | 6.327056e-08 | -7.968048e-09 | 1.003465e-09 | -1.263724e-10 |
| 3 | 1 | 0.344876 | 0.446469 | 0.288748 | 0.042585 | 0.004889 | 0.001588 | 5.391851e-05 | 1.840108e-05 | 6.279842e-06 | 2.143158e-06 | 7.314078e-07 |

▼ Training

```
1 # Initialize the weights with zero
2 theta = np.array([0.0]*len(X_train.columns))
3
4 # Initialize the weights with random values
5 theta = np.random.rand(len(X_train.columns),)
6
7 print("notice : takes approximately 3 minutes for 5k iters")
8
9 # Start the training
10 history, loss, theta = gradientDescent(X_train,
11                                     y_train,
12                                     theta,
13                                     learning_rate,
14                                     iterations,
15                                     X_valid = X_valid,
16                                     y_valid = y_valid,
17                                     loss_fn=MSE,
18                                     loss_fn_prim=MSE_prim,
19                                     decay = 0.0)
20
21
```

notice : takes approximately 3 minutes for 5k iters

100% 1000/1000 [00:46<00:00, 21.89it/s]


▼ Plotting the fitted polynomials

```
1 # Predicting using the learned weights(theta)
2 # Not used here but useful
3 y_hat = theta*X_valid
4 y_hat = np.sum(y_hat, axis=1)
5 print(MSE(X_valid, y_valid, theta))
```

0.22926660232638735

```
1 plot_curve(X_train, y_train, theta)
```


```
2 print(f"Training Loss : {RMSE(X_train, y_train, theta)}")
3 print(f"Validation Loss : {RMSE(X_valid, y_valid, theta)}\n")
```



Training Loss : 0.5101877849255999
Validation Loss : 0.4788179218934765

▼ Interactive history viewer

```
1 from ipywidgets import interact
2 import ipywidgets as widgets
3
4 @interact(iteration = widgets.IntSlider(min=0, max=iterations-1, step=100, value=0))
5 def plot_weight_history(iteration):
6     plot_curve(X_train, y_train, history["weights"][iteration])
7
8 print("Use the slider to see the algorithms progression")
```



▼ Part 4 - Plotting every Case!

Polynomial Degree :

- 5
- 8
- 10

Loss Functions :

- MSE
- RMSE
- MAE

Iterations :

- 5000
- 10000

```
1 # iterations = 1500 #TODO
2 polynomial_degrees = [5, 8, 10]
```



```

3 loss_functions = [(MSE, MSE_prim),
4                   (RMSE, RMSE_prim),
5                   (MAE, MAE_prim)]
6 fn_labels = ["MAE", "RMSE", "MSE"]
7 iterations = 2000

1  thetas = [[[],[],[]],
2            [[],[],[]],
3            [[],[],[]]]
4
5  losses = [[[],[],[]],
6            [[],[],[]],
7            [[],[],[]]]
8
9  for i,degree in enumerate(polynomial_degrees):
10     for j, (loss_fn, loss_fn_prim) in enumerate(loss_functions):
11         print(f"degree: {degree} | loss function: {fn_labels[j]}")
12
13         # preprocess data
14         X_train_copy = X_train.copy()
15         X_valid_copy = X_valid.copy()
16
17         polynomial_to_linear_regression(X_train_copy, degree)
18         polynomial_to_linear_regression(X_valid_copy, degree)
19
20         # TODO: gradient descent
21         # Initialize the weights with random values
22         theta = np.random.rand(len(X_train.columns),)
23         # Start the training
24         history, loss, theta = gradientDescent(X_train_copy,
25                                                y_train,
26                                                theta,
27                                                learning_rate,
28                                                iterations,
29                                                X_valid = X_valid_copy,
30                                                y_valid = y_valid,
31                                                loss_fn=loss_fn,
32                                                loss_fn_prim=loss_fn_prim,
33                                                decay = 0.0)
34
35         # latest iterations theta
36         thetas[i][j].append(history["weights"][-1])
37         # halfway theta
38         thetas[i][j].append(history["weights"][int(iterations/2)-1])
39
40         # Training loss
41         losses[i][j].append(loss_fn(X_train_copy, y_train, theta))
42         # Validation loss
43         losses[i][j].append(loss_fn(X_valid_copy, y_valid, theta))
44
45         print()

```

```
degree: 5 | loss function: MAE
100% 2000/2000 [00:51<00:00, 38.92it/s]

degree: 5 | loss function: RMSE
100% 2000/2000 [01:30<00:00, 22.86it/s]

degree: 5 | loss function: MSE
100% 2000/2000 [01:33<00:00, 21.39it/s]
```

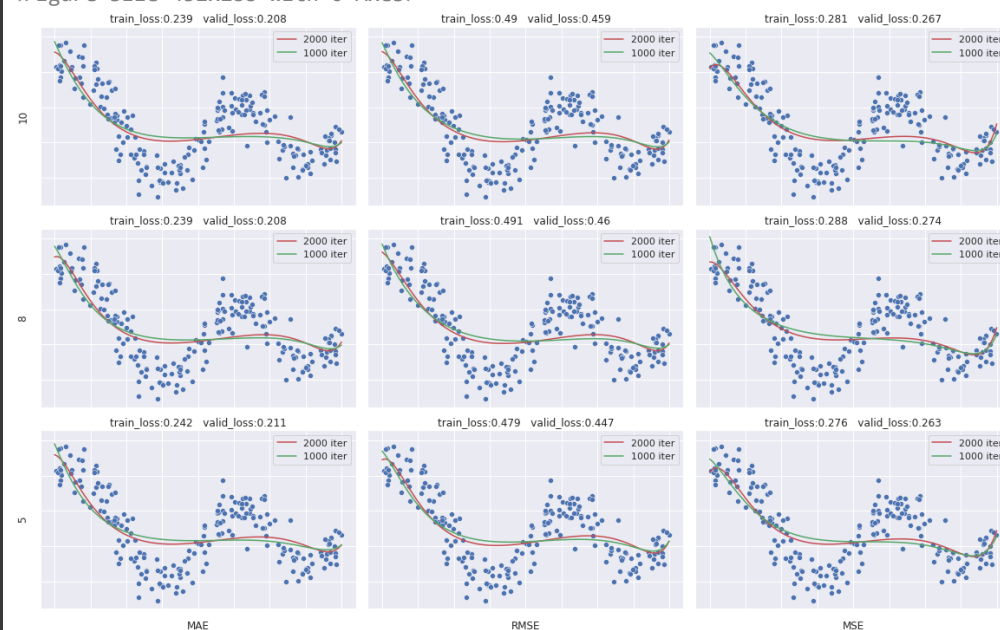
```
1 # plt.subplot_tool()
```

```
100% 2000/2000 [00:51<00:00, 38.12it/s]
```

```
1 def plot_every_curve(X, y, thetas, resolution=100):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=True, figsize=(16,10), constrained_layout=True)
8     # fig.subplots_adjust(hspace=0.6)
9
10    # Helper lists for accessing the current config
11    fn_labels = ["MAE", "RMSE", "MSE"]
12    deg_labels = [10,8,5]
13
14    # Used to plot the fitted polynomial in the range[-1,1]
15    x = np.linspace(-1, 1, resolution)
16
17    # Plotting every case in a 3 by 3 grid
18    for i in range(3):
19        for j in range(3):
20
21            # 1. plot the original data (Blue)
22            sns.scatterplot(x=X["x"], y=y, ax=axes[i,j])
23
24            # 2. Plot a curve with Last iteration theta [0] (Red)
25            theta = thetas[i][j][0]
26            sns.lineplot(x, polyCoefficients(x, theta), color='r', ax=axes[i,j])
27
28            # 3. Plot a curve with Middle iteration theta [1] (Green)
29            theta = thetas[i][j][1]
30            sns.lineplot(x, polyCoefficients(x, theta), color='g', ax=axes[i,j])
31
32            # Legends and titles
33            axes[i,j].legend(labels=[f"{iterations} iter", f"{iterations//2} iter"])
34            axes[i,j].set_title(f"train_loss:{round(losses[i][j][0],3)}    valid_loss:{round(losses[i][j][1],3)}")
35
36            # Matplotlib related code
37            axes[i,j].xaxis.set_ticklabels([])
38            axes[i,j].yaxis.set_ticklabels([])
39            axes[i,j].set_xlabel(fn_labels[j])
40            axes[i,j].set_ylabel(deg_labels[i])
41
42    plt.show()
```

```
1 plot_every_curve(X_train, y_train, thetas)
```

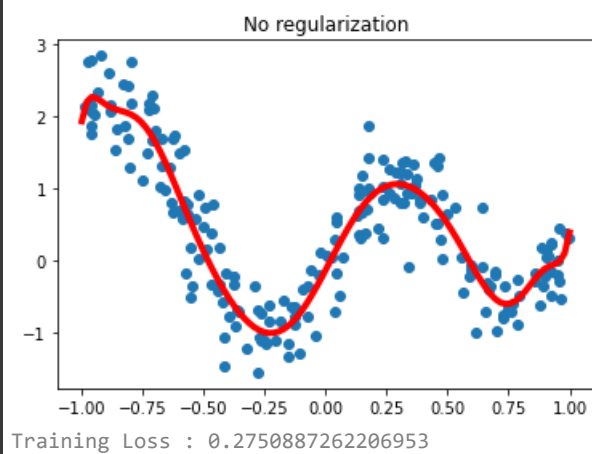
<Figure size 432x288 with 0 Axes>



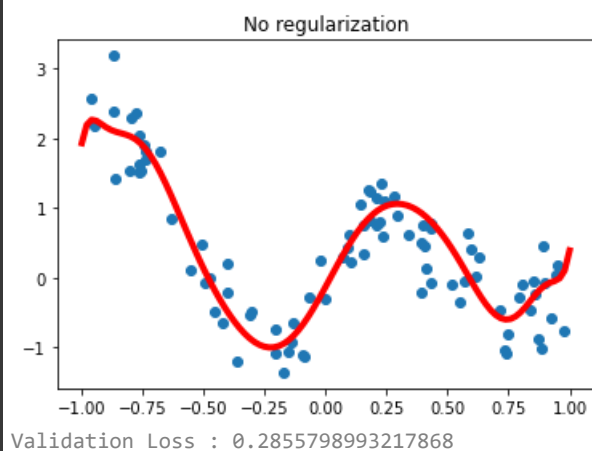
Part 5 - Normal Equation

```
1 normal_theta = normalEquation(X_train, y_train)
```

```
1 # Plotting the fitted polynomial over training data
2 plot_curve(X_train, y_train, normal_theta, title="No regularization")
3 print(f"Training Loss : {RMSE(X_train, y_train, normal_theta)}")
```




```
1 # Plotting the fitted polynomial over validation data
2 plot_curve(X_valid, y_valid, normal_theta, title="No regularization")
3 print(f"Validation Loss : {RMSE(X_valid, y_valid, normal_theta)}\n")
```




Part 6 - Regularized Normal Equation


```
1 reg_normal_theta = regularizedNormalEquation(X_train, y_train, lambda=.1)
2 plot_curve(X_train, y_train, reg_normal_theta, title="lambda=0.1")
3 print(f"Training Loss : {RMSE(X_train, y_train, reg_normal_theta)}")
4 print(f"Validation Loss : {RMSE(X_valid, y_valid, reg_normal_theta)}\n")
5
6 reg_normal_theta = regularizedNormalEquation(X_train, y_train, lambda=1)
7 plot_curve(X_train, y_train, reg_normal_theta, title="lambda=1")
8 print(f"Training Loss : {RMSE(X_train, y_train, reg_normal_theta)}")
9 print(f"Validation Loss : {RMSE(X_valid, y_valid, reg_normal_theta)}\n")
10
11 reg_normal_theta = regularizedNormalEquation(X_train, y_train, lambda=4)
12 plot_curve(X_train, y_train, reg_normal_theta, title="lambda=4")
13 print(f"Training Loss : {RMSE(X_train, y_train, reg_normal_theta)}")
14 print(f"Validation Loss : {RMSE(X_valid, y_valid, reg_normal_theta)}\n")
15
```



Training Loss : 0.34234945259970995
Validation Loss : 0.33329843645565466



Training Loss : 0.43260647240880507
Validation Loss : 0.4018612985008604



Training Loss : 0.5017630147189559
Validation Loss : 0.4717228178183581

✓ 4s completed at 4:30 AM

