

ML Exercise 1-1

Question 1

Hyper parameters

change before running the note book

learning_rate: 0.030

iterations: 10000

[Show code](#)

Loading Data

```
1 # Download dataset
2 !wget https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from tqdm.notebook import tqdm
5 import seaborn as sns
6
7 # use seaborn
8 sns.set()
9
10 # Load the data using pandas
11 df = pd.read_csv("dataset1.csv")
```

```
1 # Show a sample of the data
2 df.head()
```

	x	y
0	0.097627	0.626964
1	0.430379	0.846452
2	0.205527	0.756017
3	0.089766	0.427504
4	-0.152690	-1.335228

```
1 # Show a description of the data (might be useful later)
2 df.describe()
```

x

y

▼ Helper functions

```
std      0.589948      1.021100
```

► Loss Functions

```
[ ] ↳ 1 cell hidden
```

```
50%      0.045096      0.316442
```

▼ Gradient Descent

```
max      0.997694      3.186153
```

```
1 def gradientDescent(X, y, theta, lr, iteration, X_valid, y_valid, loss_fn = MSE, loss_fn_prim = MSE_prim, decay=0.0):
2     # Training loss per iteration history
3     train_loss_history = []
4     # Validation loss per iteration history
5     validation_loss_history = []
6     # weights progression towards the optimal value
7     theta_history = []
8
9     lr0 = lr
10    # Progress bar
11    with tqdm(total=iteration) as pbar:
12        for itera in range(iteration):
13            # TODO : Learning rate decay
14            lr = lr0 * 1/(1 + decay * itera)
15
16            for i in range(0, len(X.columns)):
17                # partial derivative of loss function with respect to Xi
18                gradient = loss_fn_prim(X, y, i, theta)
19
20                # Actual "Gradient Descent" !
21                theta[i] -= lr * gradient
22
23            # Calculating the loss after each iteration
24            # of updating the weights using Gradient Descent
25            loss = loss_fn(X, y, theta)
26            if X_valid is not None and y_valid is not None:
27                validation_loss = loss_fn(X_valid, y_valid, theta)
28
29            # Save the history of loss and weights
30            train_loss_history.append(loss)
31            if X_valid is not None and y_valid is not None:
32                validation_loss_history.append(validation_loss)
33            theta_history.append(theta.copy())
34
35            # Update progress bar
36            pbar.update(1)
37
38    history = {"training_loss":train_loss_history,
39              "validation_loss":validation_loss_history,
40              "weights":theta_history}
41    # returns loss history, latest loss, weights
42    print(f"training_loss : {round(train_loss_history[-1],4)} | validation_loss : {round(validation_loss_history[-1],4)}")
43    return history, loss, theta
```

```
1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6 def polynomial_to_linear_regression(X, polynomial_degree):
7     for i in range(2, 2 + polynomial_degree):
8         X['x'+str(i)] = X['x']**i
```

▼ Normal Equation

```
1 def normalEquation(X, y):
```

```

2 # (X^T X)^-1 X^T Y
3 XTX = np.dot(X.T,X)
4 XTX_inverse = np.linalg.inv(XTX)
5 XTY = np.dot(X.T,y)
6 theta = np.dot(XTX_inverse, XTY)
7 return theta
8
9 def regularizedNormalEquation(X, y, lambd=0.1):
10 # (X^T X + lambd I)^-1 X^T Y
11 XTX = np.dot(X.T,X) + np.dot(np.identity(X.shape[1]),lambd)
12 XTX_inverse = np.linalg.inv(XTX)
13 XTY = np.dot(X.T,y)
14 theta = np.dot(XTX_inverse, XTY)
15 return theta

```

▼ Plotting related

```

1 # helper function used to plot a polynomial
2 def polyCoefficients(x, coeffs):
3     o = len(coeffs)
4     y = 0
5     for i in range(o):
6         y += coeffs[i]*x**i
7     return y

```

```

1 # Plots a polynomial on top of the original data
2 def plot_curve(X, y, theta, c='r', title='', resolution=100):
3     plt.figure()
4     plt.title(title)
5     # Plot the original data
6     plt.scatter(x=X['x'],y= y)
7
8     x = np.linspace(-1, 1, resolution)
9     # Plot the fitted polynomial over the data
10    plt.plot(x, polyCoefficients(x, theta), c=c, linewidth=4)
11
12    plt.show()

```

```

1 def plot_every_curve(X, y, thetas, resolution=100):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=True, figsize=(16,10), constrained_layout=True)
8     # fig.subplots_adjust(hspace=0.6)
9
10    # Helper lists for accessing the current config
11    fn_labels = ["MAE", "RMSE", "MSE"]
12    deg_labels = ["degree 10","degree 8","degree 5"]
13
14    # Used to plot the fitted polynomial in the range[-1,1]
15    x = np.linspace(-1, 1, resolution)
16
17    # Plotting every case in a 3 by 3 grid
18    for i in range(3):
19        for j in range(3):
20
21            # 1. plot the original data (Blue)
22            sns.scatterplot(x=X["x"], y=y, ax=axes[i,j])
23
24            # 2. Plot a curve with Last iteration theta [0] (Red)
25            theta = thetas[i][j][0]
26            sns.lineplot(x, polyCoefficients(x, theta), color='r', ax=axes[i,j])
27
28            # 3. Plot a curve with Middle iteration theta [1] (Green)
29            theta = thetas[i][j][1]
30            sns.lineplot(x, polyCoefficients(x, theta), color='g', ax=axes[i,j])
31
32    # Legends and titles

```

```

33 axes[i,j].legend(labels=[f"{iterations} iter", f"{iterations//2} iter"])
34 axes[i,j].set_title(f"train_loss:{round(losses[i][j][0],3)}    valid_loss:{round(losses[i][j][1],3)}")
35
36 # Matplotlib related code
37 axes[i,j].xaxis.set_ticklabels([])
38 axes[i,j].yaxis.set_ticklabels([])
39 axes[i,j].set_xlabel(fn_labels[j])
40 axes[i,j].set_ylabel(deg_labels[i])
41
42 plt.show()

```

```

1 def plot_every_case_loss(histories, starting_iter=0):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=False, figsize=(16,10), constrained_layout=True)
8
9     # Helper lists for accessing the current config
10    fn_labels = ["MAE", "RMSE", "MSE"]
11    deg_labels = ["degree 10", "degree 8", "degree 5"]
12
13    # X = iterations range
14    x = np.linspace(0, iterations, iterations)
15
16    # Plotting every case in a 3 by 3 grid
17    for i in range(3):
18        for j in range(3):
19
20            # 1. Training_loss - iteration curve (Red)
21            sns.lineplot(x[starting_iter:], histories[i][j][0]["training_loss"][starting_iter:], color='r', ax=axes[i,j])
22            # 2. Validation_loss - iteration curve (green)
23            sns.lineplot(x[starting_iter:], histories[i][j][0]["validation_loss"][starting_iter:], color='g', ax=axes[i,j])
24
25            # Legends
26            axes[i,j].legend(labels=[f"training loss", f"validation loss"])
27
28            # Matplotlib related code
29            axes[i,j].set_xlabel(fn_labels[j])
30            axes[i,j].set_ylabel(deg_labels[i])
31
32    plt.show()

```


▼ Part 1 - Plotting the data

```

1 # Plot the data using matplotlib
2 df.plot(kind='scatter', x='x', y='y')
3 plt.show()

```

c argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence



▼ Part 2 - Shuffle

```

1 def plot_colorize(df):
2     '''
3     Assigns 'red' color to the first half of the data
4     and 'blue' to the rest
5
6     If the data is well shuffled we should see random red
7     and blue circles everywhere.
8
9     If the data is NOT well shuffled we might see a pattern between
10    circles' position and their color.
11    '''
12    # col = {True:"indianred", False:"cornflowerblue"}
13    col = {True:"crimson", False:"dodgerblue"}
14
15    colors = list(map(lambda x:col[x<df.shape[0]/2],df.index.tolist()))
16
17    df.plot(kind='scatter', x='x', y='y', color=colors)
18    plt.show()


```

▼ Default data

```

1 plot_colorize(df)
2 print(" "*9,"Data seems to be well shuffled.")

```




Data seems to be well shuffled.

▼ Sorted data

```

1 # Sort the data based on 'x' first, then do the
2 # previous part to see the result.
3 sorted_df = df.sort_values(by='x', ascending=True, ignore_index=True)
4
5 plot_colorize(sorted_df)
6 print(" "*6,"Data doesn't seem to be well shuffled.")

```




Data doesn't seem to be well shuffled.

▼ Shuffled data

```

1 # Let's shuffle the data anyways (just in case)
2
3 # pandas Doc: specifying drop=True prevents .reset_index()
4 # from creating a column containing the old index entries.
5 shuffled_df = df.sample(frac=1).reset_index(drop=True)
6
7 plot_colorize(shuffled_df)
8 print(" *9,\"Data seems to be well shuffled.\")

```



▼ Part 3 - Gradient Descent

Finding The optimal "Theta" values (weights)

▼ Data Prepration

▼ Add a column for bias

```

1 # Add a new column for simplicity of the calculations
2 # acts as the bias term
3 shuffled_df = pd.concat([pd.Series(1, index=shuffled_df.index, name='0'), shuffled_df], axis=1)
4 shuffled_df.head()

```

	0	x	y
0	1	0.392687	-0.202069
1	1	-0.891324	2.592220
2	1	0.957237	0.444572
3	1	0.763471	-0.603117
4	1	-0.762545	1.637844

▼ Seperate X,y

```

1 # Split training data into X and y
2 X = shuffled_df.drop(columns="y")
3 y = shuffled_df.iloc[:, 2]
4
5 print(X.head(),end="\n\n")
6 print(y.head())

```

```

0      x
0 1  0.392687
1 1 -0.891324
2 1  0.957237
3 1  0.763471
4 1 -0.762545

0  -0.202069

```

```
1      2.592220
2      0.444572
3     -0.603117
4      1.637844
Name: y, dtype: float64
```

```
1 # Split to train and valid
2 split = 0.7
3
4 X_train = X.iloc[ : int(len(X)*split),:].reset_index(drop=True)
5 X_valid = X.iloc[int(len(X)*split) : ,:].reset_index(drop=True)
6
7 y_train = y.iloc[ : int(len(X)*split)].reset_index(drop=True)
8 y_valid = y.iloc[int(len(X)*split) : ].reset_index(drop=True)
9
10 print(f"Train X size = {len(X_train)}")
11 print(f"Train y size = {len(y_train)}")
12 print(f"Valid X size = {len(X_valid)}")
13 print(f"Valid y size = {len(y_valid)}")
```

```
Train X size = 210
Train y size = 210
Valid X size = 90
Valid y size = 90
```

```
1 # Save a copy of X and y
2 # TODO might not need it
3 X_train_org = X_train.copy()
4 y_train_org = y_train.copy()
```

Polynomial Regression

a basic example

polynomial_degree:

[Show code](#)

Convert

aX + bX² + cX³ + d

to

aX1 + bX2 + cX3 + d

```
1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6 polynomial_to_linear_regression(X_train, polynomial_degree)
7 polynomial_to_linear_regression(X_valid, polynomial_degree)
8
9 X_train.head()
```

		0	x	x2	x3	x4	x5	x6	x7	x8	x9	x10	x11
0	1	0.392687	0.154203	0.060554	0.023779	0.009338	0.003667	0.001440	0.000565	0.000222	0.000087	0.000034	
1	1	-0.891324	0.794459	-0.708120	0.631164	-0.562572	0.501434	-0.446940	0.398368	-0.355075	0.316487	-0.282093	
2	1	0.957237	0.916302	0.877118	0.839609	0.803705	0.769336	0.736437	0.704944	0.674798	0.645942	0.618319	
3	1	0.763471	0.582888	0.445018	0.339758	0.259395	0.198041	0.151198	0.115435	0.088132	0.067286	0.051371	
4	1	-0.762545	0.581474	-0.443400	0.338112	-0.257826	0.196604	-0.149919	0.114320	-0.087174	0.066474	-0.050689	

Training

```

1 # Initialize the weights with zero
2 theta = np.array([0.0]*len(X_train.columns))
3
4 # Initialize the weights with random values
5 theta = np.random.rand(len(X_train.columns),)
6
7 print("notice : takes approximately 3 minutes for 5k iters")
8
9 # tip : nice way to find decay (https://machinelearningmastery.com/using-learning-rate-schedules-deep-learning-models-pyth
10 # Decay = LearningRate / Epochs
11 # Decay = 0.1 / 1000
12 # Decay = 0.0001
13
14 # Start the training
15 history, loss, theta = gradientDescent(X_train,
16                                       y_train,
17                                       theta,
18                                       learning_rate,
19                                       iterations,
20                                       X_valid = X_valid,
21                                       y_valid = y_valid,
22                                       loss_fn=MSE,
23                                       loss_fn_prim=MSE_prim,
24                                       decay = 0.0)
25
26

```

```

notice : takes approximately 3 minutes for 5k iters
100%                               10000/10000 [06:14<00:00, 26.80it/s]
training_loss : 0.1491 | validation_loss : 0.1609

```

▼ Plotting the fitted polynomials

```

1 # Predicting using the learned weights(theta)
2 # Not used here but useful
3 y_hat = theta*X_valid
4 y_hat = np.sum(y_hat, axis=1)
5 print(MSE(X_valid, y_valid, theta))


```

```
0.16093676345655536
```

```

1 plot_curve(X_train, y_train, theta)
2 print(f"Training Loss : {MSE(X_train, y_train, theta)}")
3 print(f"Validation Loss : {MSE(X_valid, y_valid, theta)}\n")

```



```


Training Loss : 0.14911607013673991
Validation Loss : 0.16093676345655536

```

```

1 plot_curve(X_train, y_train, theta)
2 print(f"Training Loss : {MSE(X_train, y_train, theta)}")
3 print(f"Validation Loss : {MSE(X_valid, y_valid, theta)}\n")


```

Training Loss : 0.14911607012672091

▼ Interactive history viewer

```
1 from ipywidgets import interact
2 import ipywidgets as widgets
3
4 @interact(iteration = widgets.IntSlider(min=0, max=iterations-1, step=100, value=0))
5 def plot_weight_history(iteration):
6     plot_curve(X_train, y_train, history["weights"][iteration])
7
8 print("Use the slider to see the algorithms progression")
```



▼ Part 4 - Plotting every Case!

Polynomial Degree :

- 5
- 8
- 10

Loss Functions :

- MSE
- RMSE
- MAE

Iterations :

- 5000
- 10000

```
1 # iterations = 100
2 polynomial_degrees = [5, 8, 10]
3 loss_functions = [(MSE, MSE_prim),
4                   (RMSE, RMSE_prim),
5                   (MAE, MAE_prim)]
6 fn_labels = ["MAE", "RMSE", "MSE"]
```

▼ Training all 9 models

```
1 # Containers used for storing results about each model
2 # Used later by the plotting functions
3 thetas = [[[],[],[]],
4           [[],[],[]],
5           [[],[],[]]]
6
7 losses = [[[],[],[]],
8           [[],[],[]],
9           [[],[],[]]]
10
11 histories = [[[],[],[]],
12             [[],[],[]],
13             [[],[],[]]]
14
15 # start the training process for each model
16 for i,degree in enumerate(polynomial_degrees):
17     for j, (loss_fn, loss_fn_prim) in enumerate(loss_functions):
18         print(f"degree: {degree} | loss function: {fn_labels[j]}")
19
20         # preprocess data (univariate non-linear to multivariate linear)
21         X_train_copy = X_train.copy()
22         X_valid_copy = X_valid.copy()
23         polynomial_to_linear_regression(X_train_copy, degree)
24         polynomial_to_linear_regression(X_valid_copy, degree)
25
26         # Initialize the weights with random values
27         theta = np.random.rand(len(X_train.columns),)
28
29         # Start the training
30         history, loss, theta = gradientDescent(X_train_copy,
31                                               y_train,
32                                               theta,
33                                               learning_rate,
34                                               iterations,
35                                               X_valid = X_valid_copy,
36                                               y_valid = y_valid,
37                                               loss_fn=loss_fn,
38                                               loss_fn_prim=loss_fn_prim,
39                                               decay = 0.0)
40
41         # Saving latest iteration's theta for each model
42         thetas[i][j].append(history["weights"][-1])
43         # Saving halfway theta for each model
44         thetas[i][j].append(history["weights"][int(iterations/2)-1])
45
46         # Saving Training loss for each model
47         losses[i][j].append(loss_fn(X_train_copy, y_train, theta))
48         # Saving Validation loss for each model
49         losses[i][j].append(loss_fn(X_valid_copy, y_valid, theta))
50         # Saving Histories for each model (used for plotting loss per iteration)
51         histories[i][j].append(history)
52
53     print()
```

```

degree: 5 | loss function: MAE
100% 10000/10000 [06:12<00:00, 27.36it/s]
training_loss : 0.1473 | validation_loss : 0.1593

degree: 5 | loss function: RMSE
100% 10000/10000 [11:12<00:00, 15.22it/s]
training_loss : 0.3774 | validation_loss : 0.3922

degree: 5 | loss function: MSE
100% 10000/10000 [06:47<00:00, 24.75it/s]
training_loss : 0.2056 | validation_loss : 0.2186

degree: 8 | loss function: MAE
100% 10000/10000 [06:12<00:00, 27.32it/s]
training_loss : 0.1475 | validation_loss : 0.1592

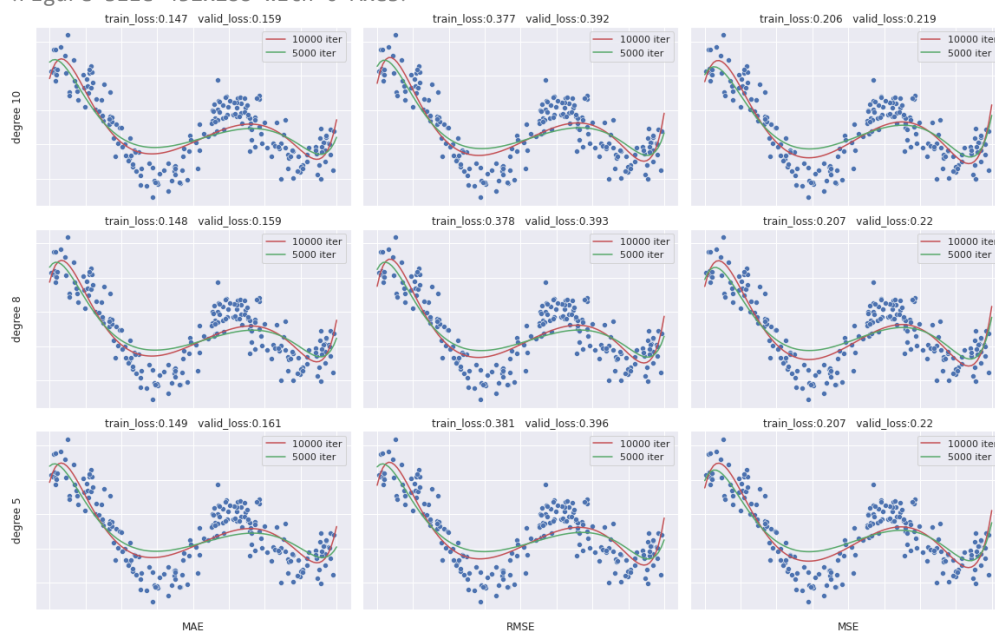
degree: 8 | loss function: RMSE
100% 10000/10000 [11:09<00:00, 14.46it/s]
training_loss : 0.3778 | validation_loss : 0.3927

degree: 8 | loss function: MSE

```

```
1 plot_every_curve(X_train, y_train, thetas)
```

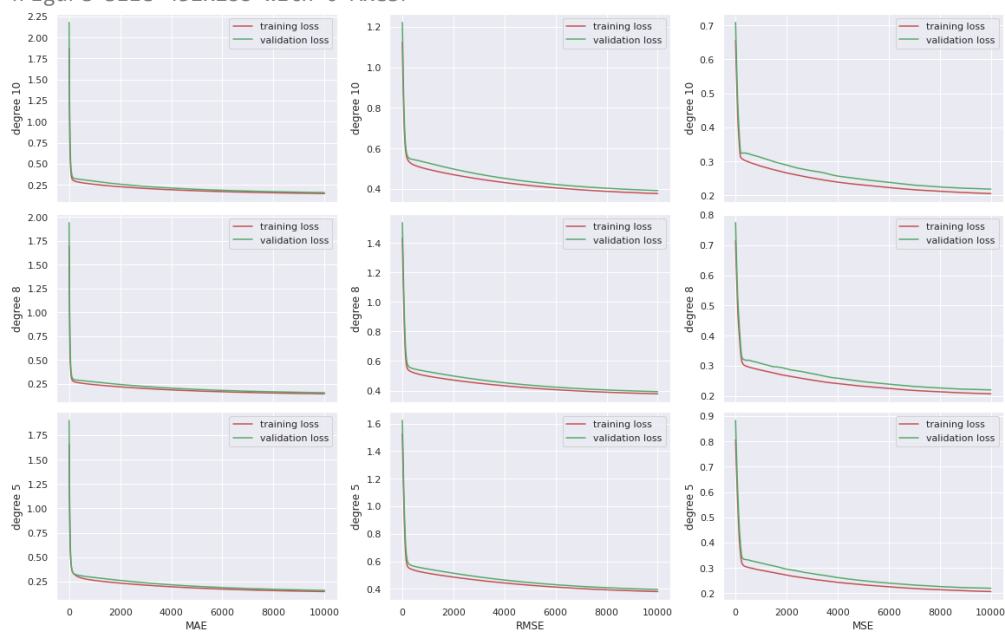
<Figure size 432x288 with 0 Axes>



Plotting train/valid loss

```
1 plot_every_case_loss(histories)
```

<Figure size 432x288 with 0 Axes>



▼ Plotting Learning Rate

```
1 _lrs = []
2 # Iterations
3 _iterations = 100
4 # Initial lr
5 _lr0 = 0.1
6 # Decay
7 _decay = _lr0/_iterations
8
9 # Simulate gradient descents main loop
10 _lr = _lr0
11 for i in range(_iterations):
12     _lr = _lr * 1/(1 + _decay * i)
13     _lrs.append(_lr)
14
15 _x = list(range(_iterations))
16 _y = _lrs
17
18 plt.plot(_x, _y)
```

```
[<matplotlib.lines.Line2D at 0x7f0646089f50>]
```

0.10

Part 5 - Normal Equation


```
1 normal_theta = normalEquation(X_train, y_train)
```

0.04

```
1 # Plotting the fitted polynomial over training data
2 plot_curve(X_train, y_train, normal_theta, title="No regularization")
3 print(f"Training Loss : {RMSE(X_train, y_train, normal_theta)}")
```




```
1 # Plotting the fitted polynomial over validation data
2 plot_curve(X_valid, y_valid, normal_theta, title="No regularization")
3 print(f"Validation Loss : {RMSE(X_valid, y_valid, normal_theta)}\n")
```



Part 6 - Regularized Normal Equation

```
1 reg_normal_theta = regularizedNormalEquation(X_train, y_train, lambda=.1)
2 plot_curve(X_train, y_train, reg_normal_theta, title="lambda=0.1")
3 print(f"Training Loss : {RMSE(X_train, y_train, reg_normal_theta)}")
4 print(f"Validation Loss : {RMSE(X_valid, y_valid, reg_normal_theta)}\n")
5
6 reg_normal_theta = regularizedNormalEquation(X_train, y_train, lambda=1)
7 plot_curve(X_train, y_train, reg_normal_theta, title="lambda=1")
8 print(f"Training Loss : {RMSE(X_train, y_train, reg_normal_theta)}")
9 print(f"Validation Loss : {RMSE(X_valid, y_valid, reg_normal_theta)}\n")
10
11 reg_normal_theta = regularizedNormalEquation(X_train, y_train, lambda=4)
12 plot_curve(X_train, y_train, reg_normal_theta, title="lambda=4")
13 print(f"Training Loss : {RMSE(X_train, y_train, reg_normal_theta)}")
14 print(f"Validation Loss : {RMSE(X_valid, y_valid, reg_normal_theta)}\n")
15
```


lambda=0.1



Training Loss : 0.34126275359191854

Validation Loss : 0.35843601252690255


lambda=1



Training Loss : 0.4250887587842727

Validation Loss : 0.44552564006365325

lambda=4



Training Loss : 0.4910232709927601

Validation Loss : 0.5202602652456745