

ML Exercise 1-1

Question 1

Hyper parameters

change before running the note book

learning_rate: 2.23

iterations: 3000

Show code

Loading Data

```
1 # Download dataset
2 !wget https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv

--2021-11-21 07:17:50-- https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.110.133, 185.199.111.133, 185.199.109.133, ..
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.110.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 12074 (12K) [text/plain]
Saving to: 'dataset1.csv.1'

dataset1.csv.1      100%[=====>] 11.79K  --.-KB/s   in 0s

2021-11-21 07:17:50 (81.5 MB/s) - 'dataset1.csv.1' saved [12074/12074]
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from tqdm.notebook import tqdm
5 import seaborn as sns
6
7 # use seaborn
8 sns.set()
9
10 # Load the data using pandas
11 df = pd.read_csv("dataset1.csv")
```

```
1 # Show a sample of the data
2 df.head()
```

	x	y
0	0.097627	0.626964
1	0.430379	0.846452
2	0.205527	0.756017
3	0.089766	0.427504
4	-0.152690	-1.335228

```
1 # Show a description of the data (might be useful later)
2 df.describe()
```

	x	y
count	300.000000	300.000000
mean	0.007005	0.412755
std	0.580948	1.021100
min	-0.990609	-1.547934
25%	-0.504657	-0.361192
50%	0.045096	0.316442
75%	0.460611	1.092441

▼ Helper functions

► Loss Functions

[] ↳ 1 cell hidden

▼ Gradient Descent

```

1 def gradientDescent(X, y, theta, lr, iteration, X_valid, y_valid, loss_fn = MSE, loss_fn_prim = MSE_prim, decay=0.0):
2     # Training loss per iteration history
3     train_loss_history = []
4     # Validation loss per iteration history
5     validation_loss_history = []
6     # weights progression towards the optimal value
7     theta_history = []
8
9     # Progress bar
10    with tqdm(total=iteration) as pbar:
11        for itera in range(iteration):
12            # TODO : Learning rate decay
13            lr = lr * 1/(1 + decay * itera)
14
15            for i in range(0, len(X.columns)):
16                # partial derivative of loss function with respect to Xi
17                gradient = loss_fn_prim(X, y, i, theta)
18
19                # Actual "Gradient Descent" !
20                theta[i] -= lr * gradient
21
22            # Calculating the loss after each iteration
23            # of updating the weights using Gradient Descent
24            loss = loss_fn(X, y, theta)
25            if X_valid is not None and y_valid is not None:
26                validation_loss = loss_fn(X_valid, y_valid, theta)
27
28            # Save the history of loss and weights
29            train_loss_history.append(loss)
30            if X_valid is not None and y_valid is not None:
31                validation_loss_history.append(validation_loss)
32            theta_history.append(theta.copy())
33
34            # Update progress bar
35            pbar.update(1)
36
37    history = {"training_loss":train_loss_history,
38              "validation_loss":validation_loss_history,
39              "weights":theta_history}
40    # returns loss history, latest loss, weights
41    print(f"training_loss : {round(train_loss_history[-1],4)} | validation_loss : {round(validation_loss_history[-1],4)}")
42    return history, loss, theta

```

```

1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression

```

```

4 # where every polynomial term is a feature for
5 # the linear regression
6 def polynomial_to_linear_regression(X, polynomial_degree):
7     for i in range(2, 2 + polynomial_degree):
8         X['x'+str(i)] = X['x']**i

```

▼ Normal Equation

```

1 def normalEquation(X, y):
2     #  $(X^T X)^{-1} X^T Y$ 
3     XTX = np.dot(X.T,X)
4     XTX_inverse = np.linalg.inv(XTX)
5     XTY = np.dot(X.T,y)
6     theta = np.dot(XTX_inverse, XTY)
7     return theta
8
9 def regularizedNormalEquation(X, y, lambd=0.1):
10    #  $(X^T X + \lambda I)^{-1} X^T Y$ 
11    XTX = np.dot(X.T,X) + np.dot(np.identity(X.shape[1]),lambd)
12    XTX_inverse = np.linalg.inv(XTX)
13    XTY = np.dot(X.T,y)
14    theta = np.dot(XTX_inverse, XTY)
15    return theta

```

▼ Plotting related

```

1 # helper function used to plot a polynomial
2 def polyCoefficients(x, coeffs):
3     o = len(coeffs)
4     y = 0
5     for i in range(o):
6         y += coeffs[i]*x**i
7     return y

```

```

1 # Plots a polynomial on top of the original data
2 def plot_curve(X, y, theta, c='r', title='', resolution=200):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
7     plt.figure()
8     plt.title(title)
9     # Plot the original data
10    plt.scatter(x=X['x'],y= y)
11
12    x = np.linspace(-1, 1, resolution)
13    # Plot the fitted polynomial over the data
14    plt.plot(x, polyCoefficients(x, theta), color=c, linewidth=4)
15
16    plt.show()

```

```

1 # Plots validation and training losses per iteration
2 def plot_loss(history, title='', starting_iter=0):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
7     fig, ax = plt.subplots()
8     plt.title(title)
9
10    # X = iterations range
11    x = np.linspace(0, iterations, iterations)
12
13    # 1. Training_loss - iteration curve (Red)
14    sns.lineplot(x[starting_iter:], history["training_loss"][starting_iter:], color='r')
15    # 2. Validation_loss - iteration curve (green)
16    sns.lineplot(x[starting_iter:], history["validation_loss"][starting_iter:], color='g')
17
18    ax.legend(labels=["training", "validation"])

```

```
19
20 plt.show()
```

```
1 def plot_lr(lr=0.1, iterations=1000, decay=None, title='learning_rate'):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     _lrs = []
7     # Iterations
8     _iterations = iterations
9     # Initial lr
10    _lr0 = lr
11    _decay = 0
12    # Decay
13    if decay is None:
14        _decay = _lr0/_iterations
15    else:
16        _decay = decay
17
18    # Simulate gradient descents main loop
19    _lr = _lr0
20    for i in range(_iterations):
21        _lr = _lr * 1/(1 + _decay * i)
22        _lrs.append(_lr)
23
24    _x = list(range(_iterations))
25    _y = _lrs
26
27    plt.figure()
28    plt.title(title)
29    plt.plot(_x, _y)
30    plt.show()
```

```
1 def plot_every_curve(X, y, thetas, resolution=100):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=True, figsize=(16,10), constrained_layout=True)
8     # fig.subplots_adjust(hspace=0.6)
9
10    # Helper lists for accessing the current config
11    fn_labels = ["MAE", "RMSE", "MSE"]
12    deg_labels = ["degree 10", "degree 8", "degree 5"]
13
14    # Used to plot the fitted polynomial in the range[-1,1]
15    x = np.linspace(-1, 1, resolution)
16
17    # Plotting every case in a 3 by 3 grid
18    for i in range(3):
19        for j in range(3):
20
21            # 1. plot the original data (Blue)
22            sns.scatterplot(x=X["x"], y=y, ax=axes[i,j])
23
24            # 2. Plot a curve with Last iteration theta [0] (Red)
25            theta = thetas[i][j][0]
26            sns.lineplot(x, polyCoefficients(x, theta), color='r', ax=axes[i,j])
27
28            # 3. Plot a curve with Middle iteration theta [1] (Green)
29            theta = thetas[i][j][1]
30            sns.lineplot(x, polyCoefficients(x, theta), color='g', ax=axes[i,j])
31
32            # Legends and titles
33            axes[i,j].legend(labels=[f"{iterations} iter", f"{iterations//2} iter"])
34            axes[i,j].set_title(f"train_loss:{round(losses[i][j][0],3)}    valid_loss:{round(losses[i][j][1],3)}")
35
36            # Matplotlib related code
37            axes[i,j].xaxis.set_ticklabels([])
```

```

38     axes[i,j].yaxis.set_ticklabels([])
39     axes[i,j].set_xlabel(fn_labels[j])
40     axes[i,j].set_ylabel(deg_labels[i])
41
42     plt.show()

```

```

1 def plot_every_case_loss(histories, starting_iter=0):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=False, figsize=(16,10), constrained_layout=True)
8
9     # Helper lists for accessing the current config
10    fn_labels = ["MAE", "RMSE", "MSE"]
11    deg_labels = ["degree 10", "degree 8", "degree 5"]
12
13    # X = iterations range
14    x = np.linspace(0, iterations, iterations)
15
16    # Plotting every case in a 3 by 3 grid
17    for i in range(3):
18        for j in range(3):
19
20            # 1. Training_loss - iteration curve (Red)
21            sns.lineplot(x[starting_iter:], histories[i][j][0]["training_loss"][starting_iter:], color='r', ax=axes[i,j])
22            # 2. Validation_loss - iteration curve (green)
23            sns.lineplot(x[starting_iter:], histories[i][j][0]["validation_loss"][starting_iter:], color='g', ax=axes[i,j])
24
25            # Legends
26            axes[i,j].legend(labels=[f"training loss", f"validation loss"])
27
28            # Matplotlib related code
29            axes[i,j].set_xlabel(fn_labels[j])
30            axes[i,j].set_ylabel(deg_labels[i])
31
32    plt.show()

```

```


1 # Plots a polynomial on top of the original data
2 def plot_normal_equations(X, y, c='r', title='', resolution=200):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
7     fig, ax = plt.subplots()
8     plt.title(title)
9     fig.dpi=120
10
11    # Plot the original data
12    sns.scatterplot(x=X['x'], y=y, size=1, color='darkgray')
13
14    x = np.linspace(-1, 1, resolution)
15
16    # calculate theta for each method
17    normal_theta = normalEquation(X_train, y_train)
18    reg_normal_theta1 = regularizedNormalEquation(X_train, y_train, lambd=0.075)
19    reg_normal_theta2 = regularizedNormalEquation(X_train, y_train, lambd=0.75)
20    reg_normal_theta4 = regularizedNormalEquation(X_train, y_train, lambd=7.5)
21
22    # Plot the fitted polynomial over the data
23    plt.plot(x, polyCoefficients(x, normal_theta), linewidth=2)
24    plt.plot(x, polyCoefficients(x, reg_normal_theta1), linewidth=2)
25    plt.plot(x, polyCoefficients(x, reg_normal_theta2), linewidth=2)
26    plt.plot(x, polyCoefficients(x, reg_normal_theta4), linewidth=2)
27
28    ax.legend(labels=["No Regularization", " $\lambda=0.075$ ", " $\lambda=0.75$ ", " $\lambda=7.5$ "])
29    ax.xaxis.set_ticklabels([])
30    ax.yaxis.set_ticklabels([])
31    ax.xaxis.set_visible(False)
32    ax.yaxis.set_visible(False)
33    plt.show()

```

▼ Part 1 - Plotting the data

```
1 # Plot the data using matplotlib
2 df.plot(kind='scatter', x='x', y='y')
3 plt.show()
```

c argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence over more descriptive names



▼ Part 2 - Shuffle

```
1 def plot_colorize(df):
2     '''
3     Assigns 'red' color to the first half of the data
4     and 'blue' to the rest
5
6     If the data is well shuffled we should see random red
7     and blue circles everywhere.
8
9     If the data is NOT well shuffled we might see a pattern between
10    circles' position and their color.
11    '''
12
13    df_red = df.loc[df.index<df.shape[0]/2]
14    df_blue = df.loc[df.index>=df.shape[0]/2]
15
16    sns.scatterplot(data=df_red, x='x', y='y', color='r')
17    sns.scatterplot(data=df_blue, x='x', y='y', color='b')
18
19    plt.show()
```

▼ Default data

```
1 plot_colorize(df)
2 print(" *9,\"Data seems to be well shuffled.\")
```

▼ Sorted data

```
1 # Sort the data based on 'x' first, then do the
2 # previous part to see the result.
3 sorted_df = df.sort_values(by='x', ascending=True, ignore_index=True)
4
5 plot_colorize(sorted_df)
6 print(" "*6, "Data doesn't seem to be well shuffled.")
```



▼ Shuffled data

```
1 # Let's shuffle the data anyways (just in case)
2
3 # pandas Doc: specifying drop=True prevents .reset_index()
4 # from creating a column containing the old index entries.
5 shuffled_df = df.sample(frac=1).reset_index(drop=True)
6
7 plot_colorize(shuffled_df)
8 print(" "*9, "Data seems to be well shuffled.")
```



▼ Part 3 - Gradient Descent

Finding The optimal "Theta" values (weights)

▼ Data Prepration

▼ Add a column for bias

```
1 # Add a new column for simplicity of the calculations
2 # acts as the bias term
3 shuffled_df = pd.concat([pd.Series(1, index=shuffled_df.index, name='0'), shuffled_df], axis=1)
4 shuffled_df.head()
```

	0	x	y
0	1	-0.582246	0.725637
1	1	-0.726199	2.085352
2	1	0.561058	-0.445023
3	1	0.530651	0.739235
4	1	0.395262	0.760835

▼ Seperate X,y

```
1 # Split training data into X and y
2 X = shuffled_df.drop(columns="y")
3 y = shuffled_df.iloc[:, 2]
4
5 print(X.head(),end="\n\n")
6 print(y.head())
```

```

0      x
0  1 -0.582246
1  1 -0.726199
2  1  0.561058
3  1  0.530651
4  1  0.395262
```

```

0      0.725637
1      2.085352
2     -0.445023
3      0.739235
4      0.760835
Name: y, dtype: float64
```

```
1 # Split to train and valid
2 split = 0.7
3
4 X_train = X.iloc[ : int(len(X)*split),:].reset_index(drop=True)
5 X_valid = X.iloc[int(len(X)*split) : ,:].reset_index(drop=True)
6
7 y_train = y.iloc[ : int(len(X)*split)].reset_index(drop=True)
8 y_valid = y.iloc[int(len(X)*split) : ].reset_index(drop=True)
9
10 print(f"Train X size = {len(X_train)}")
11 print(f"Train y size = {len(y_train)}")
12 print(f"Valid X size = {len(X_valid)}")
13 print(f"Valid y size = {len(y_valid)}")
```

```

Train X size = 210
Train y size = 210
Valid X size = 90
Valid y size = 90
```

```
1 # Save a copy of X and y
2 # TODO might not need it
3 X_train_org = X_train.copy()
4 y_train_org = y_train.copy()
```

▼ Polynomial Regression

a basic example

polynomial_degree:

[Show code](#)

Convert

to

```
1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6 polynomial_to_linear_regression(X_train, polynomial_degree)
7 polynomial_to_linear_regression(X_valid, polynomial_degree)
8
9 X_train.head()
```

		0	x	x2	x3	x4	x5	x6	x7	x8	x9	x10	x11
0	1		-0.582246	0.339011	-0.197388	0.114928	-0.066917	0.038962	-0.022685	0.013209	-0.007691	0.004478	-0.002607
1	1		-0.726199	0.527366	-0.382973	0.278115	-0.201967	0.146668	-0.106510	0.077348	-0.056170	0.040791	-0.029622
2	1		0.561058	0.314786	0.176614	0.099091	0.055596	0.031192	0.017501	0.009819	0.005509	0.003091	0.001734
3	1		0.530651	0.281590	0.149426	0.079293	0.042077	0.022328	0.011848	0.006287	0.003336	0.001770	0.000939
4	1		0.395262	0.156232	0.061753	0.024409	0.009648	0.003813	0.001507	0.000596	0.000235	0.000093	0.000037

▼ Training

Normal equation for comparison :

- 0.07 training loss
- 0.09 validation loss

5k iter 10th degree polynomial $lr=2.3$

- Training Loss : 0.07766969752936674
- Validation Loss : 0.094857479611961

[illegible]

notice : takes approximately 3 minutes for 5k iters

100% 3000/3000 [01:42<00:00, 28.61it/s]


training_loss : 0.0974 | validation_loss : 0.0851

▼ Plotting the fitted polynomials

```
1 # Predicting using the learned weights(theta)
2 # Not used here but useful
3 y_hat = theta*X_valid
4 y_hat = np.sum(y_hat, axis=1)
5 print(MSE(X_valid, y_valid, theta))
```


0.0851166629303032

```
1 plot_curve(X_train, y_train, theta, title='Linear Regression')
2
3 print(f"Training Loss : {MSE(X_train, y_train, theta)}")
4 print(f"Validation Loss : {MSE(X_valid, y_valid, theta)}\n")
```



▼ Plotting losses

```
1 plot_loss(history, starting_iter=10, title='Losses')
```




▼ Plotting the learning rate

```
1 plot_lr(lr=learning_rate, iterations=iterations, decay=decay)
```



▼ Interactive history viewer

```
1 from ipywidgets import interact
2 import ipywidgets as widgets
3
4 @interact(iteration = widgets.IntSlider(min=0, max=iterations-1, step=100, value=0))
5 def plot_weight_history(iteration):
6     plot_curve(X_train, y_train, history["weights"][iteration])
7
8 print("Use the slider to see the algorithms progression")
```



▼ Part 4 - Plotting every Case!

Polynomial Degree :

- 5
- 8
- 10

Loss Functions :

- MSE
- RMSE
- MAE

Iterations :

- 5000
- 10000

```
1 # iterations = 100
2 polynomial_degrees = [5, 8, 10]
3 loss_functions = [(MSE, MSE_prim),
4                   (RMSE, RMSE_prim),
5                   (MAE, MAE_prim)]
6 fn_labels = ["MAE", "RMSE", "MSE"]
```

▼ Training all 9 models

```
1 # Containers used for storing results about each model
2 # Used later by the plotting functions
3 thetas = [[[]], [], []],
```

```

4         [[],[],[]],
5         [[],[],[]]]
6
7 losses = [[[],[],[]],
8           [[],[],[]],
9           [[],[],[]]]
10
11 histories = [[[],[],[]],
12             [[],[],[]],
13             [[],[],[]]]
14
15 # start the training process for each model
16 for i,degree in enumerate(polynomial_degrees):
17     for j, (loss_fn, loss_fn_prim) in enumerate(loss_functions):
18         print(f"degree: {degree} | loss function: {fn_labels[j]}")
19
20         # preprocess data (univariate non-linear to multivariate linear)
21         X_train_copy = X_train.copy()
22         X_valid_copy = X_valid.copy()
23         polynomial_to_linear_regression(X_train_copy, degree)
24         polynomial_to_linear_regression(X_valid_copy, degree)
25
26         # Initialize the weights with random values
27         theta = np.random.rand(len(X_train.columns),)
28
29         # Start the training
30         history, loss, theta = gradientDescent(X_train_copy,
31                                               y_train,
32                                               theta,
33                                               learning_rate,
34                                               iterations,
35                                               X_valid = X_valid_copy,
36                                               y_valid = y_valid,
37                                               loss_fn=loss_fn,
38                                               loss_fn_prim=loss_fn_prim,
39                                               decay = 0.0)
40
41         # Saving latest iteration's theta for each model
42         thetas[i][j].append(history["weights"][-1])
43         # Saving halfway theta for each model
44         thetas[i][j].append(history["weights"][int(iterations/2)-1])
45
46         # Saving Training loss for each model
47         losses[i][j].append(loss_fn(X_train_copy, y_train, theta))
48         # Saving Validation loss for each model
49         losses[i][j].append(loss_fn(X_valid_copy, y_valid, theta))
50         # Saving Histories for each model (used for plotting loss per iteration)
51         histories[i][j].append(history)
52
53     print()

```

```

degree: 5 | loss function: MAE
100%                                     3000/3000 [01:43<00:00, 29.71it/s]
training_loss : 0.0957 | validation_loss : 0.083

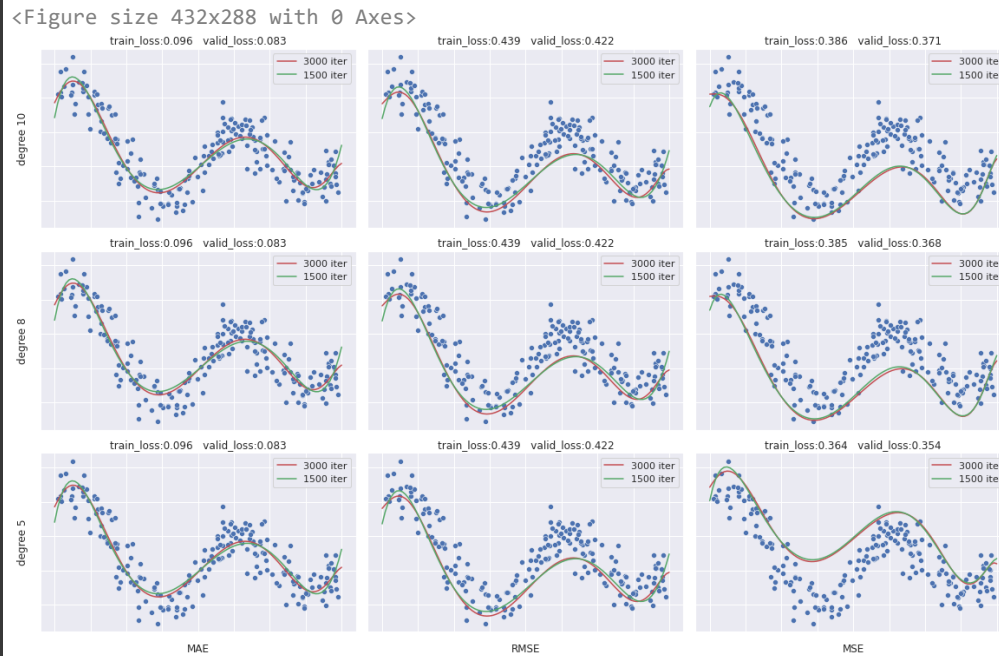
degree: 5 | loss function: RMSE
100%                                     3000/3000 [03:05<00:00, 16.16it/s]
training_loss : 0.4388 | validation_loss : 0.422

degree: 5 | loss function: MSE
100%                                     3000/3000 [01:53<00:00, 26.02it/s]
training_loss : 0.3856 | validation_loss : 0.3707

degree: 8 | loss function: MAE
100%                                     3000/3000 [01:41<00:00, 30.32it/s]
training_loss : 0.0959 | validation_loss : 0.0832

```

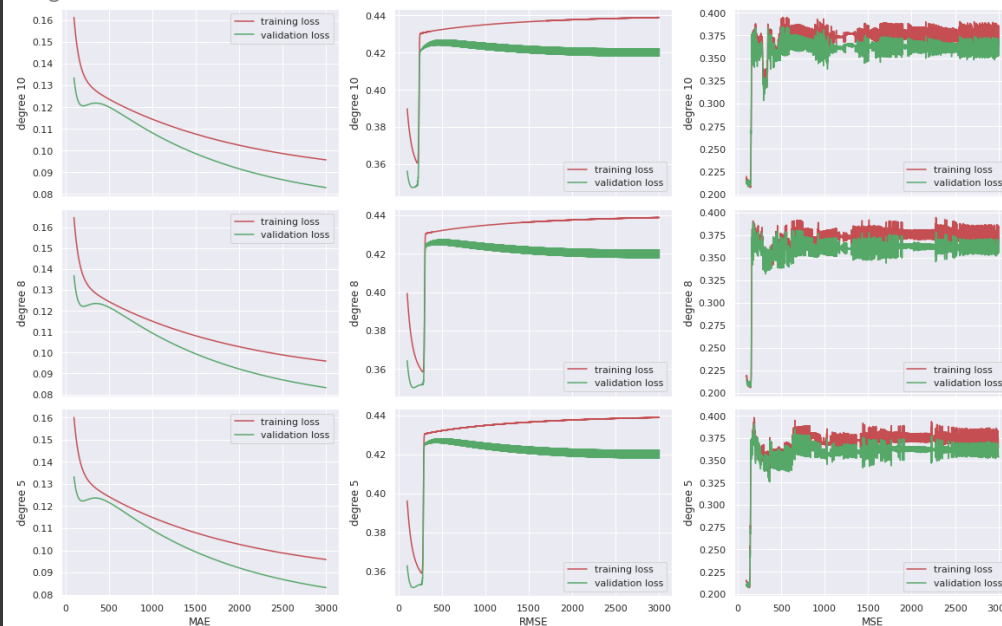
```
1 plot_every_curve(X_train, y_train, thetas)
```



Plotting train/valid loss


```
1 plot_every_case_loss(histories, starting_iter=100)
```

<Figure size 432x288 with 0 Axes>



▼ Plotting Learning Rate

```
1 plot_lr(lr=learning_rate, iterations=iterations, decay=decay)
```




▼ Part 5 - Normal Equation

```
1 normal_theta = normalEquation(X_train, y_train)
```


```
1 # Plotting the fitted polynomial over training data
2 plot_curve(X_train, y_train, normal_theta, title="No regularization")
3 print(f"Training Loss : {RMSE(X_train, y_train, normal_theta)}")
```

No regularization



```
1 # Plotting the fitted polynomial over validation data
2 plot_curve(X_valid, y_valid, normal_theta, title="No regularization")
3 print(f"Validation Loss : {RMSE(X_valid, y_valid, normal_theta)}\n")
```

No regularization



Validation Loss : 0.28624132493466486

▼ Part 6 - Regularized Normal Equation

```
1 loss_fn = RMSE # as asked in the question
2
3 reg_normal_theta1 = regularizedNormalEquation(X_train, y_train, lambd=0.075)
4 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=0.1")
5 reg_normal_theta2 = regularizedNormalEquation(X_train, y_train, lambd=0.75)
6 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=1")
7 reg_normal_theta3 = regularizedNormalEquation(X_train, y_train, lambd=7.5)
8 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=4")
```

```
1 names = ["λ = 0.075", "λ = 0.75", "λ = 7.5"]
2
3 training_errors = [round(loss_fn(X_train, y_train, reg_normal_theta1),5),
4                    round(loss_fn(X_train, y_train, reg_normal_theta2),5),
5                    round(loss_fn(X_train, y_train, reg_normal_theta3),5)]
6
7 validation_errors =[round(loss_fn(X_valid, y_valid, reg_normal_theta1),5),
8                    round(loss_fn(X_valid, y_valid, reg_normal_theta2),5),
9                    round(loss_fn(X_valid, y_valid, reg_normal_theta3),5)]
10
11 for i in range(len(names)):
12     print(f"{i} - Regularized Normal Equation ({names[i]})")
13     print(f"Training Loss : {training_errors[i]}")
14     print(f"Validation Loss : {validation_errors[i]}\n")
```

0 - Regularized Normal Equation ($\lambda = 0.075$)
 Training Loss : 0.34547
 Validation Loss : 0.32648

1 - Regularized Normal Equation ($\lambda = 0.75$)
 Training Loss : 0.42381
 Validation Loss : 0.38665

2 - Regularized Normal Equation ($\lambda = 7.5$)
 Training Loss : 0.52674
 Validation Loss : 0.49811

```
1 plot_normal_equations(X_train, y_train)
```

