

ML Exercise 1-1

Question 1

Hyper parameters

change before running the note book

learning_rate: 2.23

iterations: 5000

Show code

Loading Data

```
1 # Download dataset
2 !wget https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv

--2021-11-21 09:16:49-- https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.110.133, 185.199.111.133, 185.199.108.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.110.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 12074 (12K) [text/plain]
Saving to: 'dataset1.csv.11'

dataset1.csv.11      100%[=====>] 11.79K  --.-KB/s   in 0s

2021-11-21 09:16:49 (110 MB/s) - 'dataset1.csv.11' saved [12074/12074]
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from tqdm.notebook import tqdm
5 import seaborn as sns
6
7 # use seaborn
8 sns.set()
9
10 # Load the data using pandas
11 df = pd.read_csv("dataset1.csv")
```

```
1 # Show a sample of the data
2 df.head()
```

	x	y
0	0.097627	0.626964
1	0.430379	0.846452
2	0.205527	0.756017
3	0.089766	0.427504
4	-0.152690	-1.335228

```
1 # Show a description of the data (might be useful later)
2 df.describe()
```

	x	y
count	300.000000	300.000000
mean	0.007005	0.412755
std	0.580948	1.021100
min	-0.990609	-1.547934
25%	-0.504657	-0.361192
50%	0.045096	0.316442
75%	0.460611	1.092441

▼ Helper functions

▼ Loss Functions

```

1 # h(theta) = theta transpose * X
2 def hypothesis(X, theta):
3     y1 = theta*X
4     return np.sum(y1, axis=1)
5
6 def MSE(X, y, theta):
7     y_hat = hypothesis(X, theta)
8     m = len(X)
9     return sum((y_hat-y)**2)/(2*m)
10
11 def RMSE(X, y, theta):
12     y_hat = hypothesis(X, theta)
13     m = len(X)
14     return np.sqrt(sum((y_hat-y)**2)/(2*m))
15
16 def MAE(X, y, theta):
17     y_hat = hypothesis(X, theta)
18     m = len(X)
19     return sum(np.abs((y_hat-y)))/(2*m)
20
21 # Loss functions Derivatives
22 def MSE_prim(X, y, i, theta):
23     y_hat = hypothesis(X, theta)
24     Xi = X.iloc[:, i]
25     m = len(X)
26     return sum((y_hat-y) * Xi) / m
27
28 def RMSE_prim(X, y, i, theta):
29     # src : https://math.stackexchange.com/questions/4065532/rmse-derivatives
30     mse = MSE(X, y, theta)
31     mse_prim = MSE_prim(X, y, i, theta)
32
33     return mse_prim / 2 / np.sqrt(mse)
34
35 def MAE_prim(X, y, i, theta):
36     # src : https://stats.stackexchange.com/questions/312737/mean-absolute-error-mae-derivative
37     # src2 : https://github.com/chenxingwei/machine_learning_from_scratch/blob/master/algorithm/2.linearRegressionGradientDescent.py
38     y_hat = hypothesis(X, theta)
39     # print(np.sum((X.T*(np.sign(y_hat-y)/len(X))), axis=1)[i])
40     return np.sum((X.T*(np.sign(y_hat-y)/len(X))), axis=1)[i]
41

```

▼ Gradient Descent

```

1 def gradientDescent(X, y, theta, lr, iteration, X_valid, y_valid, loss_fn = MSE, loss_fn_prim = MSE_prim, decay=0.0):
2     # Training loss per iteration history
3     train_loss_history = []
4     # Validation loss per iteration history
5     validation_loss_history = []
6     # weights progression towards the optimal value

```

```

7  theta_history = []
8
9  # Progress bar
10 with tqdm(total=iteration) as pbar:
11     for itera in range(iteration):
12         # TODO : Learning rate decay
13         lr = lr * 1/(1 + decay * itera)
14
15         for i in range(0, len(X.columns)):
16             # partial derivative of loss function with respect to Xi
17             gradient = loss_fn_prim(X, y, i, theta)
18
19             # Actual "Gradient Descent" !
20             theta[i] -= lr * gradient
21
22         # Calculating the loss after each iteration
23         # of updating the weights using Gradient Descent
24         loss = loss_fn(X, y, theta)
25         if X_valid is not None and y_valid is not None:
26             validation_loss = loss_fn(X_valid, y_valid, theta)
27
28         # Save the history of loss and weights
29         train_loss_history.append(loss)
30         if X_valid is not None and y_valid is not None:
31             validation_loss_history.append(validation_loss)
32         theta_history.append(theta.copy())
33
34         # Update progress bar
35         pbar.update(1)
36
37 history = {"training_loss":train_loss_history,
38           "validation_loss":validation_loss_history,
39           "weights":theta_history}
40 # returns loss history, latest loss, weights
41 print(f"training_loss : {round(train_loss_history[-1],4)} | validation_loss : {round(validation_loss_history[-1],4)}")
42 return history, loss, theta

```

```

1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6 def polynomial_to_linear_regression(X, polynomial_degree):
7     for i in range(2, 1 + polynomial_degree):
8         X['x'+str(i)] = X['x']**i

```

▼ Normal Equation

```

1 def normalEquation(X, y):
2     #  $(X^T X)^{-1} X^T Y$ 
3     XTX = np.dot(X.T,X)
4     XTX_inverse = np.linalg.inv(XTX)
5     XTY = np.dot(X.T,y)
6     theta = np.dot(XTX_inverse, XTY)
7     return theta
8
9 def regularizedNormalEquation(X, y, lambd=0.1):
10    #  $(X^T X + \lambda I)^{-1} X^T Y$ 
11    XTX = np.dot(X.T,X) + np.dot(np.identity(X.shape[1]),lambd)
12    XTX_inverse = np.linalg.inv(XTX)
13    XTY = np.dot(X.T,y)
14    theta = np.dot(XTX_inverse, XTY)
15    return theta

```

▼ Plotting related

```

1 # helper function used to plot a polynomial
2 def polyCoefficients(x, coeffs):
3     o = len(coeffs)

```

```

4     y = 0
5     for i in range(o):
6         y += coeffs[i]*x**i
7     return y

```

```

1 # Plots a polynomial on top of the original data
2 def plot_curve(X, y, theta, c='r', title='', resolution=200):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
7     plt.figure()
8     plt.title(title)
9     # Plot the original data
10    plt.scatter(x=X['x'],y= y)
11
12    x = np.linspace(-1, 1, resolution)
13    # Plot the fitted polynomial over the data
14    plt.plot(x, polyCoefficients(x, theta), color=c, linewidth=4)
15
16    plt.show()

```

```

1 # Plots validation and training losses per iteration
2 def plot_loss(history, title='', starting_iter=0):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
7     fig, ax = plt.subplots()
8     plt.title(title)
9
10    # X = iterations range
11    x = np.linspace(0, iterations, iterations)
12
13    # 1. Training_loss - iteration curve (Red)
14    sns.lineplot(x[starting_iter:], history["training_loss"][starting_iter:], color='r')
15    # 2. Validation_loss - iteration curve (green)
16    sns.lineplot(x[starting_iter:], history["validation_loss"][starting_iter:], color='g')
17
18    ax.legend(labels=["training", "validation"])
19
20    plt.show()

```

```

1 def plot_lr(lr=0.1, iterations=1000, decay=None, title='learning_rate'):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     _lrs = []
7     # Iterations
8     _iterations = iterations
9     # Initial lr
10    _lr0 = lr
11    _decay = 0
12    # Decay
13    if decay is None:
14        _decay = _lr0/_iterations
15    else:
16        _decay = decay
17
18    # Simulate gradient descents main loop
19    _lr = _lr0
20    for i in range(_iterations):
21        _lr = _lr * 1/(1 + _decay * i)
22        _lrs.append(_lr)
23
24    _x = list(range(_iterations))
25    _y = _lrs
26
27    plt.figure()
28    plt.title(title)

```

```

29 plt.plot(_x, _y)
30 plt.show()

1 def plot_every_curve(X, y, thetas, resolution=100):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=True, figsize=(16,10), constrained_layout=True)
8     # fig.subplots_adjust(hspace=0.6)
9
10    # Helper lists for accessing the current config
11    fn_labels = ["MAE", "RMSE", "MSE"]
12    deg_labels = ["degree 10", "degree 8", "degree 3"]
13
14    # Used to plot the fitted polynomial in the range[-1,1]
15    x = np.linspace(-1, 1, resolution)
16
17    # Plotting every case in a 3 by 3 grid
18    for i in range(3):
19        for j in range(3):
20
21            # 1. plot the original data (Blue)
22            sns.scatterplot(x=X["x"], y=y, ax=axes[i,j])
23
24            # 2. Plot a curve with Last iteration theta [0] (Red)
25            theta = thetas[i][j][0]
26            sns.lineplot(x, polyCoefficients(x, theta), color='r', ax=axes[i,j])
27
28            # 3. Plot a curve with Middle iteration theta [1] (Green)
29            theta = thetas[i][j][1]
30            sns.lineplot(x, polyCoefficients(x, theta), color='g', ax=axes[i,j])
31
32            # Legends and titles
33            axes[i,j].legend(labels=[f"{iterations} iter", f"{iterations//2} iter"])
34            axes[i,j].set_title(f"train_loss:{round(losses[i][j][0],3)}   valid_loss:{round(losses[i][j][1],3)}")
35
36            # Matplotlib related code
37            axes[i,j].xaxis.set_ticklabels([])
38            axes[i,j].yaxis.set_ticklabels([])
39            axes[i,j].set_xlabel(fn_labels[j])
40            axes[i,j].set_ylabel(deg_labels[i])
41
42    plt.show()

```

```

1 def plot_every_case_loss(histories, starting_iter=0):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=False, figsize=(16,10), constrained_layout=True)
8
9     # Helper lists for accessing the current config
10    fn_labels = ["MAE", "RMSE", "MSE"]
11    deg_labels = ["degree 10", "degree 8", "degree 3"]
12
13    # X = iterations range
14    x = np.linspace(0, iterations, iterations)
15
16    # Plotting every case in a 3 by 3 grid
17    for i in range(3):
18        for j in range(3):
19
20            # 1. Training_loss - iteration curve (Red)
21            sns.lineplot(x[starting_iter:], histories[i][j][0]["training_loss"][starting_iter:], color='r', ax=axes[i,j])
22            # 2. Validation_loss - iteration curve (green)
23            sns.lineplot(x[starting_iter:], histories[i][j][0]["validation_loss"][starting_iter:], color='g', ax=axes[i,j])
24
25            # Legends
26            axes[i,j].legend(labels=[f"training loss", f"validation loss"])
27

```

```

28 # Matplotlib related code
29 axes[i,j].set_xlabel(fn_labels[j])
30 axes[i,j].set_ylabel(deg_labels[i])
31
32 plt.show()

1 # Plots a polynomial on top of the original data
2 def plot_normal_equations(X, y, c='r', title='', resolution=200):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
7     fig, ax = plt.subplots()
8     plt.title(title)
9     fig.dpi=120
10
11 # Plot the original data
12 sns.scatterplot(x=X['x'], y=y, size=1, color='darkgray')
13
14 x = np.linspace(-1, 1, resolution)
15
16 # calculate theta for each method
17 normal_theta = normalEquation(X_train_copy, y_train)
18 reg_normal_theta1 = regularizedNormalEquation(X_train_copy, y_train, lambda=0.075)
19 reg_normal_theta2 = regularizedNormalEquation(X_train_copy, y_train, lambda=0.75)
20 reg_normal_theta4 = regularizedNormalEquation(X_train_copy, y_train, lambda=7.5)
21
22 # Plot the fitted polynomial over the data
23 plt.plot(x, polyCoefficients(x, normal_theta), linewidth=2)
24 plt.plot(x, polyCoefficients(x, reg_normal_theta1), linewidth=2)
25 plt.plot(x, polyCoefficients(x, reg_normal_theta2), linewidth=2)
26 plt.plot(x, polyCoefficients(x, reg_normal_theta4), linewidth=2)
27
28 ax.legend(labels=["No Regularization", " $\lambda=0.075$ ", " $\lambda=0.75$ ", " $\lambda=7.5$ "])
29 ax.xaxis.set_ticklabels([])
30 ax.yaxis.set_ticklabels([])
31 ax.xaxis.set_visible(False)
32 ax.yaxis.set_visible(False)
33 plt.show()

```


▼ Part 1 - Plotting the data

```

1 # Plot the data using matplotlib
2 df.plot(kind='scatter', x='x', y='y')
3 plt.show()

```

c argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence



▼ Part 2 - Shuffle

```

1 def plot_colorize(df):
2     '''
3     Assigns 'red' color to the first half of the data

```

```

4 and 'blue' to the rest
5
6 If the data is well shuffled we should see random red
7 and blue circles everywhere.
8
9 If the data is NOT well shuffled we might see a pattern between
10 circles' position and their color.
11 '''
12
13 df_red = df.loc[df.index<df.shape[0]/2]
14 df_blue = df.loc[df.index>=df.shape[0]/2]
15
16 sns.scatterplot(data=df_red, x='x', y='y', color='r')
17 sns.scatterplot(data=df_blue, x='x', y='y', color='b')
18
19 plt.show()

```

▼ Default data

```

1 plot_colorize(df)
2 print(" *9,"Data seems to be well shuffled.")

```



▼ Sorted data

```

1 # Sort the data based on 'x' first, then do the
2 # previous part to see the result.
3 sorted_df = df.sort_values(by='x', ascending=True, ignore_index=True)
4
5 plot_colorize(sorted_df)
6 print(" *6,"Data doesn't seem to be well shuffled.")

```



▼ Shuffled data

```


1 # Let's shuffle the data anyways (just in case)
2

```

```

3 # pandas Doc: specifying drop=True prevents .reset_index()
4 # from creating a column containing the old index entries.
5 shuffled_df = df.sample(frac=1).reset_index(drop=True)
6
7 plot_colorize(shuffled_df)
8 print(" *9,\"Data seems to be well shuffled.\")

```



Data seems to be well shuffled.

▼ Part 3 - Gradient Descent

Finding The optimal "Theta" values (weights)

▼ Data Prepration

▼ Add a column for bias

```

1 #Add a new column for simplicity of the calculations
2 #acts as the bias term
3 shuffled_df = pd.concat([pd.Series(1, index=shuffled_df.index, name='0'), shuffled_df], axis=1)
4 shuffled_df.head()

```

	0	x	y
0	1	-0.244496	-1.151692
1	1	0.152315	0.344465
2	1	-0.492117	-0.081580
3	1	0.355633	0.936222
4	1	-0.628728	0.676142

▼ Seperate X,y

```

1 #Split training data into X and y
2 X = shuffled_df.drop(columns="y")
3 y = shuffled_df.iloc[:, 2]
4
5 print(X.head(), end="\n\n")
6 print(y.head())

```

	0	x
0	1	-0.244496
1	1	0.152315
2	1	-0.492117
3	1	0.355633
4	1	-0.628728

0	-1.151692
1	0.344465
2	-0.081580
3	0.936222


```
4      0.676142
      Name: y, dtype: float64
```

```
1 # Split to train and valid
2 split = 0.7
3
4 X_train = X.iloc[ : int(len(X)*split),:].reset_index(drop=True)
5 X_valid = X.iloc[int(len(X)*split) : ,:].reset_index(drop=True)
6
7 y_train = y.iloc[ : int(len(X)*split)].reset_index(drop=True)
8 y_valid = y.iloc[int(len(X)*split) : ].reset_index(drop=True)
9
10 print(f"Train X size = {len(X_train)}")
11 print(f"Train y size = {len(y_train)}")
12 print(f"Valid X size = {len(X_valid)}")
13 print(f"Valid y size = {len(y_valid)}")
```

```
Train X size = 210
Train y size = 210
Valid X size = 90
Valid y size = 90
```

```
1 # Save a copy of X and y
2 # TODO might not need it
3 X_train_org = X_train.copy()
4 y_train_org = y_train.copy()
```

Polynomial Regression

a basic example

polynomial_degree: 10

[Show code](#)

Convert

aX + bX² + cX³ + d

to

aX1 + bX2 + cX3 + d

```
1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6
7 X_train_copy = X_train.copy()
8 X_valid_copy = X_valid.copy()
9 polynomial_to_linear_regression(X_train_copy, polynomial_degree)
10 polynomial_to_linear_regression(X_valid_copy, polynomial_degree)
11
12
13 X_train_copy.head()
```

	0	x	x2	x3	x4	x5	x6	x7	x8	x9	x10
0	1	-0.244496	0.059778	-0.014616	0.003573	-0.000874	0.000214	-0.000052	1.276964e-05	-3.122130e-06	7.633493e-07
1	1	0.152315	0.023200	0.003534	0.000538	0.000082	0.000012	0.000002	2.896902e-07	4.412407e-08	6.720743e-09
2	1	-0.492117	0.242179	-0.119180	0.058651	-0.028863	0.014204	-0.006990	3.439893e-03	-1.692829e-03	8.330694e-04
3	1	0.355633	0.126475	0.044979	0.015996	0.005689	0.002023	0.000719	2.558687e-04	9.099537e-05	3.236096e-05
4	1	-0.628728	0.395299	-0.248536	0.156261	-0.098246	0.061770	-0.038837	2.441760e-02	-1.535203e-02	9.652255e-03

Training

Normal equation for comparison :

- 0.07 training loss
- 0.09 validation loss

5k iter 10th degree polynomial lr=2.3

- Training Loss : 0.07766969752936674
- Validation Loss : 0.094857479611961

```
1 # Initialize the weights with zero
2 theta = np.array([0.0]*len(X_train_copy.columns))
3
4 # Initialize the weights with random values
5 theta = np.random.rand(len(X_train_copy.columns),)
6
7 print("notice : takes approximately 3 minutes for 5k iters")
8
9 # tip : nice way to find decay (https://machinelearningmastery.com/using-learning-rate-schedules-deep-learning-models-pyth
10 # Decay = LearningRate / Epochs
11 # Decay = 0.1 / 1000
12 # Decay = 0.0001
13 decay = learning_rate/iterations/11250
14
15
16 # For testing purposes
17 # iterations = 500
18 # learning_rate = 2.3
19 # decay = learning_rate/iterations/250
20
21 # Start the training
22 history, loss, theta = gradientDescent(X_train_copy,
23                                       y_train,
24                                       theta,
25                                       learning_rate,
26                                       iterations,
27                                       X_valid = X_valid_copy,
28                                       y_valid = y_valid,
29                                       loss_fn=MSE,
30                                       loss_fn_prim=MSE_prim,
31                                       decay = 0)
32
33
```

notice : takes approximately 3 minutes for 5k iters

100% 5000/5000 [02:31<00:00, 33.06it/s]

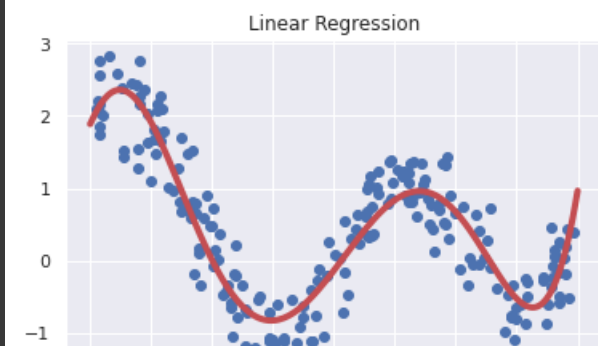
training_loss : 0.0807 | validation_loss : 0.1239

▼ Plotting the fitted polynomials

```
1 # Predicting using the learned weights(theta)
2 # Not used here but useful
3 y_hat = theta*X_valid_copy
4 y_hat = np.sum(y_hat, axis=1)
5 print(MSE(X_valid_copy, y_valid, theta))
```

0.12387766832808893


```
1 plot_curve(X_train_copy, y_train, theta, title='Linear Regression')
2
3 print(f"Training Loss : {MSE(X_train_copy, y_train, theta)}")
4 print(f"Validation Loss : {MSE(X_valid_copy, y_valid, theta)}\n")
```



▼ Plotting losses


Training Loss : 0.08069536588067559

```
1 plot_loss(history, starting_iter=0, title='Losses')
```




▼ Plotting the learning rate

```
1 plot_lr(lr=learning_rate, iterations=iterations, decay=decay)
```



▼ Interactive history viewer

```
1 from ipywidgets import interact
2 import ipywidgets as widgets
3
4 @interact(iteration = widgets.IntSlider(min=0, max=iterations-1, step=100, value=0))
5 def plot_weight_history(iteration):
6     plot_curve(X_train_copy, y_train, history["weights"][iteration])
7
8 print("Use the slider to see the algorithms progression")
```



▼ Part 4 - Plotting every Case!

Polynomial Degree :

- 5
- 8
- 10

Loss Functions :

- MSE
- RMSE
- MAE

Iterations :

- 5000
- 10000

```
1 # iterations = 500
2 polynomial_degrees = [10, 8, 3]
3 loss_functions = [(MAE, MAE_prim),
4                  (RMSE, RMSE_prim),
5                  (MSE, MSE_prim)]
6 fn_labels = ["MAE", "RMSE", "MSE"]
```

▼ Training all 9 models

```
1 # Containers used for storing results about each model
2 # Used later by the plotting functions
3 thetas = [[[],[],[]],
4           [[[],[],[]],
5           [[[],[],[]]]
6
7 losses = [[[],[],[]],
8           [[[],[],[]],
9           [[[],[],[]]]
10
11 histories = [[[],[],[]],
12              [[[],[],[]],
13              [[[],[],[]]]
14
15 # start the training process for each model
16 for i,degree in enumerate(polynomial_degrees):
17     for j, (loss_fn, loss_fn_prim) in enumerate(loss_functions):
18         print(f"degree: {degree} | loss function: {fn_labels[j]}")
19
20     # preprocess data (univariate non-linear to multivariate linear)
21     X_train_copy = X_train.copy()
22     X_valid_copy = X_valid.copy()
23     polynomial_to_linear_regression(X_train_copy, degree)
24     polynomial_to_linear_regression(X_valid_copy, degree)
25
26     # Initialize the weights with random values
27     theta = np.random.rand(len(X_train_copy.columns),)
28
29     # Start the training
30     history, loss, theta = gradientDescent(X_train_copy,
31                                           y_train,
32                                           theta,
```

```

33         learning_rate,
34         iterations,
35         X_valid = X_valid_copy,
36         y_valid = y_valid,
37         loss_fn=loss_fn,
38         loss_fn_prim=loss_fn_prim,
39         decay = decay*0)
40
41     # Saving latest iteration's theta for each model
42     thetas[i][j].append(history["weights"][-1])
43     # Saving halfway theta for each model
44     thetas[i][j].append(history["weights"][int(iterations/2)-1])
45
46     # Saving Training loss for each model
47     losses[i][j].append(loss_fn(X_train_copy, y_train, theta))
48     # Saving Validation loss for each model
49     losses[i][j].append(loss_fn(X_valid_copy, y_valid, theta))
50     # Saving Histories for each model (used for plotting loss per iteration)
51     histories[i][j].append(history)
52
53     print()

```

```

degree: 10 | loss function: MAE
100%                    5000/5000 [02:52<00:00, 28.87it/s]
training_loss : 0.3667 | validation_loss : 0.4209

degree: 10 | loss function: RMSE
100%                    5000/5000 [04:33<00:00, 17.76it/s]
training_loss : 0.431 | validation_loss : 0.4617

degree: 10 | loss function: MSE
100%                    5000/5000 [02:30<00:00, 33.36it/s]
training_loss : 0.0806 | validation_loss : 0.1237

degree: 8 | loss function: MAE
100%                    5000/5000 [02:24<00:00, 34.66it/s]
training_loss : 0.3503 | validation_loss : 0.3544

degree: 8 | loss function: RMSE
100%                    5000/5000 [03:31<00:00, 22.96it/s]
training_loss : 0.4342 | validation_loss : 0.5045

degree: 8 | loss function: MSE
100%                    5000/5000 [01:58<00:00, 42.39it/s]
training_loss : 0.0858 | validation_loss : 0.1384

degree: 3 | loss function: MAE
100%                    5000/5000 [01:11<00:00, 70.39it/s]
training_loss : 0.2812 | validation_loss : 0.2957

degree: 3 | loss function: RMSE
100%                    5000/5000 [01:18<00:00, 65.90it/s]
training_loss : 0.4708 | validation_loss : 0.4771

degree: 3 | loss function: MSE
100%                    5000/5000 [00:47<00:00, 104.60it/s]
training_loss : 0.2216 | validation_loss : 0.2277

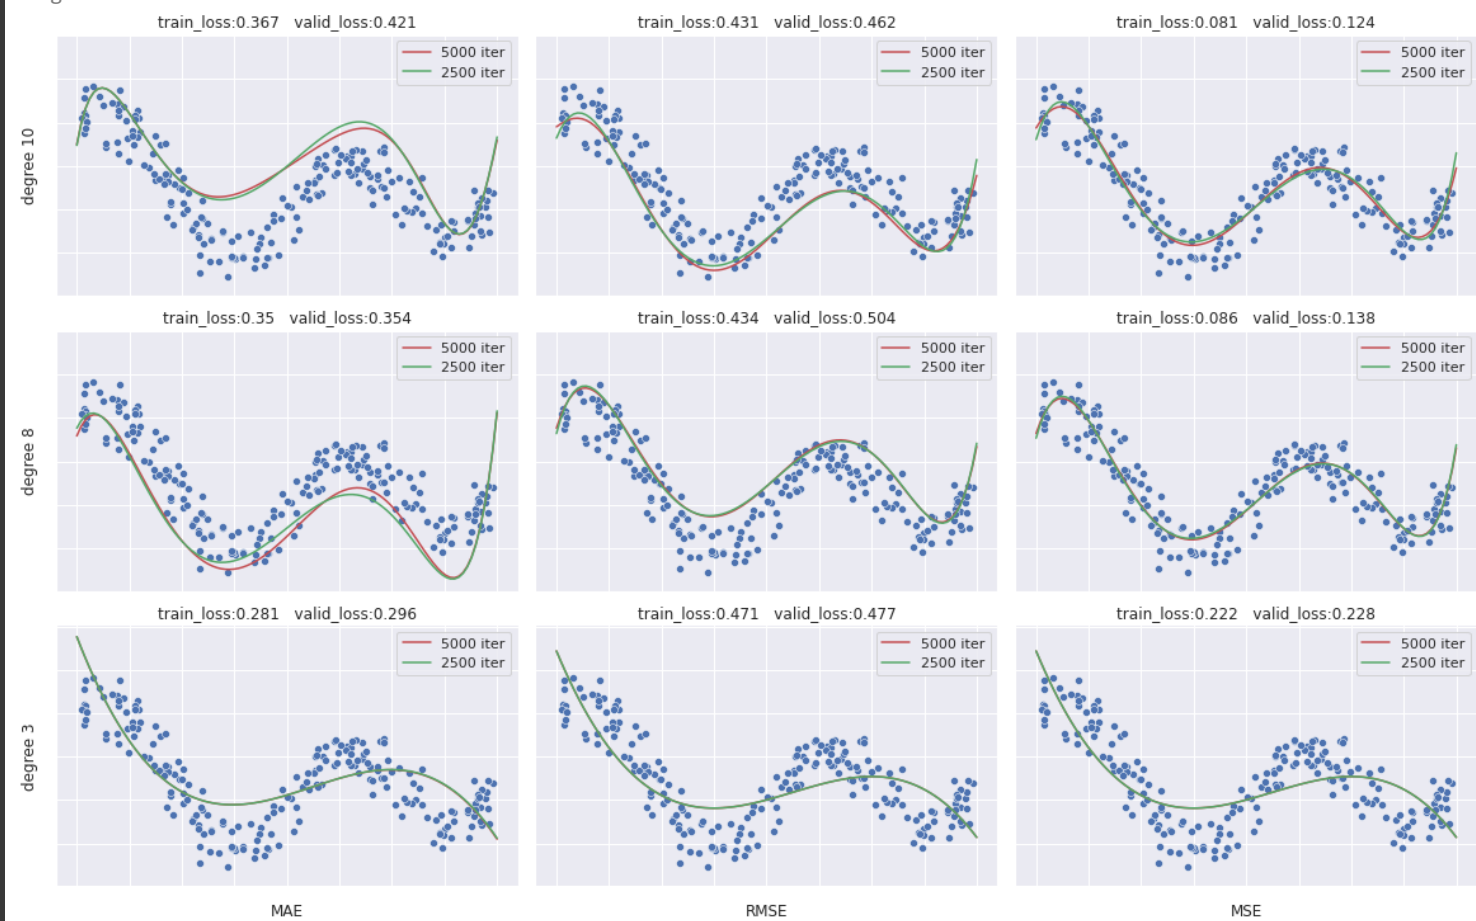
```

```

1 plot_every_curve(X_train, y_train, thetas)

```


<Figure size 432x288 with 0 Axes>




Plotting train/valid loss

```
1 plot_every_case_loss(histories, starting_iter=0)
```


<Figure size 432x288 with 0 Axes>



Plotting Learning Rate



```
1 plot_lr(lr=learning_rate, iterations=iterations, decay=decay)
```




Part 5 - Normal Equation

```
1 # preprocess data (univariate non-linear to multivariate linear)
2 X_train_copy = X_train.copy()
3 X_valid_copy = X_valid.copy()
4 polynomial_to_linear_regression(X_train_copy, 8)
5 polynomial_to_linear_regression(X_valid_copy, 8)
```

```
1 normal_theta = normalEquation(X_train_copy, y_train)
```


```
1 # Plotting the fitted polynomial over training data
2 plot_curve(X_train_copy, y_train, normal_theta, title="No regularization")
3 print(f"Training Loss : {RMSE(X_train_copy, y_train, normal_theta)}")
```



Training Loss : 0.2654493437506085

```
1 # Plotting the fitted polynomial over validation data
2 plot_curve(X_valid_copy, y_valid, normal_theta, title="No regularization")
3 print(f"Validation Loss : {RMSE(X_valid_copy, y_valid, normal_theta)}\n")
```

No regularization



Part 6 - Regularized Normal Equation

```

1 loss_fn = RMSE # as asked in the question
2
3 reg_normal_theta1 = regularizedNormalEquation(X_train_copy, y_train, lambda=0.075)
4 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=0.1")
5 reg_normal_theta2 = regularizedNormalEquation(X_train_copy, y_train, lambda=0.75)
6 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=1")
7 reg_normal_theta3 = regularizedNormalEquation(X_train_copy, y_train, lambda=7.5)
8 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=4")

```

```

1 names = ["λ = 0.075", "λ = 0.75", "λ = 7.5"]
2
3 training_errors = [round(loss_fn(X_train_copy, y_train, reg_normal_theta1),5),
4                    round(loss_fn(X_train_copy, y_train, reg_normal_theta2),5),
5                    round(loss_fn(X_train_copy, y_train, reg_normal_theta3),5)]
6
7 validation_errors =[round(loss_fn(X_valid_copy, y_valid, reg_normal_theta1),5),
8                    round(loss_fn(X_valid_copy, y_valid, reg_normal_theta2),5),
9                    round(loss_fn(X_valid_copy, y_valid, reg_normal_theta3),5)]
10
11 for i in range(len(names)):
12     print(f"{i} - Regularized Normal Equation ({names[i]})")
13     print(f"Training Loss : {training_errors[i]}")
14     print(f"Validation Loss : {validation_errors[i]}\n")

```

```

0 - Regularized Normal Equation (λ = 0.075)
Training Loss : 0.31678
Validation Loss : 0.36842

```

```

1 - Regularized Normal Equation (λ = 0.75)
Training Loss : 0.43112
Validation Loss : 0.43171

```

```


2 - Regularized Normal Equation (λ = 7.5)
Training Loss : 0.52495
Validation Loss : 0.49915

```

```

1 plot_normal_equations(X_train, y_train)

```



✓ 0s completed at 1:20 PM

