

▼ ML Exercise 1-1

Question 1

By Gholamreza Dar

<https://gholamrezadar.ir/>

Nov 2021

Questions available at : [Github Link](#)

How to Run: Runtime > Run all(ctrl+f9)

▼ Hyper Parameters

Hyper Parameters

Main :

max_iterations: 10000

train_test_split: 0.7

Basic Example :

basic_polynomial_degree: 10

basic_iterations: 1000

basic_learning_rate: 2.3

basic_decay: 5e-9

Show code

▼ Loading Data

```
1 # Download dataset
2 !wget https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv

--2021-11-23 05:46:33-- https://raw.githubusercontent.com/Gholamrezadar/machine-learning-exercises/main/dataset1/dataset1.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 12074 (12K) [text/plain]
Saving to: 'dataset1.csv.1'

dataset1.csv.1      100%[=====>] 11.79K  --.-KB/s   in 0s

2021-11-23 05:46:33 (83.2 MB/s) - 'dataset1.csv.1' saved [12074/12074]
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from tqdm.notebook import tqdm
5 import seaborn as sns
6
7 # use seaborn
8 sns.set()
9
10 # Load the data using pandas
11 df = pd.read_csv("dataset1.csv")
```

```
1 # Show a sample of the data
2 df.head()
```

	x	y
0	0.097627	0.626964
1	0.430379	0.846452
2	0.205527	0.756017
3	0.089766	0.427504
4	-0.152690	-1.335228

```
1 # Show a description of the data (might be useful later)
2 df.describe()
```

	x	y
count	300.000000	300.000000
mean	0.007005	0.412755
std	0.580948	1.021100
min	-0.990609	-1.547934
25%	-0.504657	-0.361192
50%	0.045096	0.316442
75%	0.460611	1.092441
max	0.997694	3.186153

▼ Helper functions

```
1 # Suppress some warnings
2 import warnings
3 warnings.simplefilter(action="ignore", category=FutureWarning)
```

▼ Loss Functions

```
1 # h(theta) = theta transpose * X
2 def hypothesis(X, theta):
3     y1 = theta*X
4     return np.sum(y1, axis=1)
5
6 def MSE(X, y, theta):
7     y_hat = hypothesis(X, theta)
8     m = len(X)
9     return sum((y_hat-y)**2)/(2*m)
10
11 def RMSE(X, y, theta):
12     y_hat = hypothesis(X, theta)
13     m = len(X)
14     return np.sqrt(sum((y_hat-y)**2)/(2*m))
15
16 def MAE(X, y, theta):
17     y_hat = hypothesis(X, theta)
18     m = len(X)
19     return sum(np.abs((y_hat-y)))/(2*m)
```

```
20
21 # Loss functions Derivatives
22 def MSE_prim(X, y, i, theta):
23     y_hat = hypothesis(X, theta)
24     Xi = X.iloc[:, i]
25     m = len(X)
26     return sum((y_hat-y) * Xi) / m
27
28 def RMSE_prim(X, y, i, theta):
29     # src : https://math.stackexchange.com/questions/4065532/rmse-derivatives
30     mse = MSE(X, y, theta)
31     mse_prim = MSE_prim(X, y, i, theta)
32
33     return mse_prim / 2 / np.sqrt(mse)
34
35 def MAE_prim(X, y, i, theta):
36     # src : https://stats.stackexchange.com/questions/312737/mean-absolute-error-mae-derivative
37     # src2 : https://github.com/chenxingwei/machine_learning_from_scratch/blob/master/algorithm/2.LinearRegressionGradientDescent.md
38     y_hat = hypothesis(X, theta)
39     # print(np.sum((X.T*(np.sign(y_hat-y)/len(X))), axis=1)[i])
40     return np.sum((X.T*(np.sign(y_hat-y)/len(X))), axis=1)[i]
41
```

▼ Gradient Descent

```
1 def gradientDescent(X, y, theta, lr, iteration, X_valid, y_valid, loss_fn = MSE, loss_fn_prim = MSE_prim, decay=0.0):
2     # Training loss per iteration history
3     train_loss_history = []
4     # Validation loss per iteration history
5     validation_loss_history = []
6     # weights progression towards the optimal value
7     theta_history = []
8
9     # Progress bar
10    with tqdm(total=iteration) as pbar:
11        for itera in range(iteration):
12            # TODO : Learning rate decay
13            lr = lr * 1/(1 + decay * itera)
14
15            for i in range(0, len(X.columns)):
16                # partial derivative of loss function with respect to Xi
17                gradient = loss_fn_prim(X, y, i, theta)
18
19                # Actual "Gradient Descent" !
20                theta[i] -= lr * gradient
21
22            # Calculating the loss after each iteration
23            # of updating the weights using Gradient Descent
24            loss = loss_fn(X, y, theta)
25            if X_valid is not None and y_valid is not None:
26                validation_loss = loss_fn(X_valid, y_valid, theta)
27
28            # Save the history of loss and weights
29            train_loss_history.append(loss)
30            if X_valid is not None and y_valid is not None:
31                validation_loss_history.append(validation_loss)
32            theta_history.append(theta.copy())
33
34            # Update progress bar
35            pbar.update(1)
36
37    history = {"training_loss":train_loss_history,
38              "validation_loss":validation_loss_history,
39              "weights":theta_history}
40    # returns loss history, latest loss, weights
41    print(f"training_loss : {round(train_loss_history[-1],4)} | validation_loss : {round(validation_loss_history[-1],4)}")
42    return history, loss, theta
```

```
1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6 def polynomial_to_linear_regression(X, polynomial_degree):
7     for i in range(2, 1 + polynomial_degree):
8         X['x'+str(i)] = X['x']**i
```

```
1 # helper function used to plot a polynomial
2 def polyCoefficients(x, coeffs):
3     o = len(coeffs)
4     y = 0
5     for i in range(o):
6         y += coeffs[i]*x**i
7     return y
```

▼ Normal Equation

```
1 def normalEquation(X, y):
2     # (X^T X)^-1 X^T Y
3     XTX = np.dot(X.T,X)
4     XTX_inverse = np.linalg.inv(XTX)
5     XTY = np.dot(X.T,y)
6     theta = np.dot(XTX_inverse, XTY)
7     return theta
8
9 def regularizedNormalEquation(X, y, lambd=0.1):
10    # (X^T X + lambda I)^-1 X^T Y
11    XTX = np.dot(X.T,X) + np.dot(np.identity(X.shape[1]),lambd)
12    XTX_inverse = np.linalg.inv(XTX)
13    XTY = np.dot(X.T,y)
14    theta = np.dot(XTX_inverse, XTY)
15    return theta
```

▼ Plotting related

▼ def plot_curve()

```
1 # Plots a polynomial on top of the original data
2 def plot_curve(X, y, theta, c='r', title='', resolution=200):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
7     plt.figure()
8     plt.title(title)
9     # Plot the original data
10    plt.scatter(x=X['x'],y= y)
11
12    x = np.linspace(-1, 1, resolution)
13    # Plot the fitted polynomial over the data
14    plt.plot(x, polyCoefficients(x, theta), color=c, linewidth=4)
15
16    plt.show()
```

▼ def plot_loss()

```
1 # Plots validation and training losses per iteration
2 def plot_loss(history, title='', starting_iter=0):
3     # Don't try this at home
4     import warnings
5     warnings.simplefilter(action="ignore", category=FutureWarning)
6
```

```
7 fig, ax = plt.subplots()
8 plt.title(title)
9
10 # X = iterations range
11 x = np.linspace(0, iterations, iterations)
12
13 # 1. Training_loss - iteration curve (Red)
14 sns.lineplot(x[starting_iter:], history["training_loss"][starting_iter:], color='r')
15 # 2. Validation_loss - iteration curve (green)
16 sns.lineplot(x[starting_iter:], history["validation_loss"][starting_iter:], color='g')
17
18 ax.legend(labels=["training", "validation"])
19
20 plt.show()
```

▼ def plot_lr()

```
1 def plot_lr(lr=0.1, iterations=1000, decay=None, title='learning_rate'):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     _lrs = []
7     # Iterations
8     _iterations = iterations
9     # Initial lr
10    _lr0 = lr
11    _decay = 0
12    # Decay
13    if decay is None:
14        _decay = _lr0/_iterations
15    else:
16        _decay = decay
17
18    # Simulate gradient descents main loop
19    _lr = _lr0
20    for i in range(_iterations):
21        _lr = _lr * 1/(1 + _decay * i)
22        _lrs.append(_lr)
23
24    _x = list(range(_iterations))
25    _y = _lrs
26
27    plt.figure()
28    plt.title(title)
29    plt.plot(_x, _y)
30    plt.show()
```

▼ def plot_every_curve()

```
1 def plot_every_curve(X, y, thetas,line_widths=[1, 1], resolution=100, i_max=3, j_max=3, dpi=72):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=True, figsize=(16,10), constrained_layout=True)
8     fig.dpi = dpi
9     # fig.subplots_adjust(hspace=0.6)
10
11    # Helper lists for accessing the current config
12    fn_labels = ["MAE", "RMSE", "MSE"]
13    deg_labels = ["degree 10","degree 8","degree 5"]
14
15    # Used to plot the fitted polynomial in the range[-1,1]
16    x = np.linspace(-1, 1, resolution)
17
18    # Plotting every case in a 3 by 3 grid
19    for i in range(i_max):
20        for j in range(j_max):
21
22            # 1. plot the original data (Dots)
23            sns.scatterplot(x=X["x"], y=y, ax=axes[i,j] )# ,s=70, color='darkgray')
24
25            # 2. Plot a curve with an earlier iteration theta [1] (Red)
26            theta = thetas[i][j][1]
27            sns.lineplot(x, polyCoefficients(x, theta), color='r', ax=axes[i,j], linewidth=line_widths[1])
28
29            # 3. Plot a curve with the Last iteration theta [0] (Green)
30            theta = thetas[i][j][0]
31            sns.lineplot(x, polyCoefficients(x, theta), color='g', ax=axes[i,j], linewidth=line_widths[0])
32
33            # Legends and titles
34            axes[i,j].legend(labels=[f"{other_iteration_to_display} iters", f"{int(iterations)} iters"])
35            axes[i,j].set_title(f"train_loss:{round(losses[i][j][0],3)}    valid_loss:{round(losses[i][j][1],3)}")
36
37            # Matplotlib related code
38            axes[i,j].xaxis.set_ticklabels([])
39            axes[i,j].yaxis.set_ticklabels([])
40            axes[i,j].set_xlabel(fn_labels[j])
41            axes[i,j].set_ylabel(deg_labels[i])
42
43    plt.show()
```

▼ def plot_every_case_loss()

```
1 def plot_every_case_loss(histories, starting_iter=0, i_max=3, j_max=3, dpi=72):
2     # Don't try this at home
3     import warnings
4     warnings.simplefilter(action="ignore", category=FutureWarning)
5
6     plt.figure()
7     fig, axes = plt.subplots(3, 3, sharex=True, sharey=False, figsize=(16,10), constrained_layout=True)
8     fig.dpi = dpi
9
10    # Helper lists for accessing the current config
11    fn_labels = ["MAE", "RMSE", "MSE"]
12    deg_labels = ["degree 10","degree 8","degree 3"]
13
14    # X = iterations range
15    x = np.linspace(0, iterations, iterations)
16
17    # Plotting every case in a 3 by 3 grid
18    for i in range(i_max):
19        for j in range(j_max):
20
21            # 1. Training_loss - iteration curve (Red)
22            sns.lineplot(x[starting_iter:], histories[i][j][0]["training_loss"][starting_iter:], color='r', ax=axes[i,j])
23            # 2. Validation_loss - iteration curve (green)
24            sns.lineplot(x[starting_iter:], histories[i][j][0]["validation_loss"][starting_iter:], color='g', ax=axes[i,j])
25
26            # Legends
27            axes[i,j].legend(labels=[f"training loss", f"validation loss"])
28
29            # Matplotlib related code
30            axes[i,j].set_xlabel(fn_labels[j])
31            axes[i,j].set_ylabel(deg_labels[i])
32
33    plt.show()
```


▼ def plot_normal_equations()

```
1 # Plots a polynomial on top of the original data
2 def plot_normal_equations(X, y, c='r', title='', resolution=200, dpi=72):
3     # Don't try this at home
```

```
4 import warnings
5 warnings.simplefilter(action="ignore", category=FutureWarning)
6
7 fig, ax = plt.subplots()
8 plt.title(title)
9 fig.dpi=dpi
10
11 # Plot the original data
12 sns.scatterplot(x=X['x'], y=y, size=1, color='darkgray')
13
14 x = np.linspace(-1, 1, resolution)
15
16 # calculate theta for each method
17 normal_theta = normalEquation(X_train_copy, y_train)
18 reg_normal_theta1 = regularizedNormalEquation(X_train_copy, y_train, lambda=0.075)
19 reg_normal_theta2 = regularizedNormalEquation(X_train_copy, y_train, lambda=0.75)
20 reg_normal_theta4 = regularizedNormalEquation(X_train_copy, y_train, lambda=7.5)
21
22 # Plot the fitted polynomial over the data
23 plt.plot(x, polyCoefficients(x, normal_theta), linewidth=2)
24 plt.plot(x, polyCoefficients(x, reg_normal_theta1), linewidth=2)
25 plt.plot(x, polyCoefficients(x, reg_normal_theta2), linewidth=2)
26 plt.plot(x, polyCoefficients(x, reg_normal_theta4), linewidth=2)
27
28 ax.legend(labels=["No Regularization", " $\lambda=0.075$ ", " $\lambda=0.75$ ", " $\lambda=7.5$ "])
29 ax.xaxis.set_ticklabels([])
30 ax.yaxis.set_ticklabels([])
31 ax.xaxis.set_visible(False)
32 ax.yaxis.set_visible(False)
33 plt.show()
```

Part 1 - Plotting the data

```
1 # Plot the data using matplotlib
2 plt.scatter(x=df['x'],y=df['y'])
3 plt.show()
```



Part 2 - Shuffle

```
1 def plot_colorize(df):
2     '''
3     Assigns 'red' color to the first half of the data
4     and 'blue' to the rest
5
6     If the data is well shuffled we should see random red
7     and blue circles everywhere.
8
9     If the data is NOT well shuffled we might see a pattern between
10    circles' position and their color.
11    '''
12
13    df_red = df.loc[df.index<df.shape[0]/2]
14    df_blue = df.loc[df.index>=df.shape[0]/2]
15
16    sns.scatterplot(data=df_red, x='x', y='y', color='r')
17    sns.scatterplot(data=df_blue, x='x', y='y', color='b')
18
19    plt.show()
```

Default data

```
1 plot_colorize(df)
2 print(" *9,"Data seems to be well shuffled.")
```




Sorted data

```
1 # Sort the data based on 'x' first, then do the
2 # previous part to see the result.
3 sorted_df = df.sort_values(by='x', ascending=True, ignore_index=True)
4
5 plot_colorize(sorted_df)
6 print(" *6,"Data doesn't seem to be well shuffled.")
```



Shuffled data

```
1 # Let's shuffle the data anyways (just in case)
2
3 # pandas Doc: specifying drop=True prevents .reset_index()
4 # from creating a column containing the old index entries.
5 shuffled_df = df.sample(frac=1).reset_index(drop=True)
6
7 plot_colorize(shuffled_df)
8 print(" *9,"Data seems to be well shuffled.")
```



Part 3 - Gradient Descent

Finding The optimal "Theta" values (weights)

Data Preperation

Add a column for bias

```
1 # Add a new column for simplicity of the calculations
2 # acts as the bias term
3 shuffled_df = pd.concat([pd.Series(1, index=shuffled_df.index, name='0'), shuffled_df], axis=1)
4 shuffled_df.head()
```

	0	x	y
0	1	-0.122797	-0.700428
1	1	0.954990	-0.287307
2	1	-0.087699	-1.126259
3	1	-0.405126	0.206917
4	1	0.154457	0.754608

Seperate X,y

```
1 # Split training data into X and y
2 X = shuffled_df.drop(columns="y")
3 y = shuffled_df.iloc[:, 2]
4
5 print("X : \n", X.head().to_string(),end="\n\n")
6 print("y : \n", y.head().to_string())
```

X :		
0	1	x
0	1	-0.122797
1	1	0.954990
2	1	-0.087699
3	1	-0.405126
4	1	0.154457
y :		
0		-0.700428
1		-0.287307
2		-1.126259
3		0.206917
4		0.754608

```
1 # Split to train and valid
2 split = train_test_split
3
4 X_train = X.iloc[ : int(len(X)*split),:].reset_index(drop=True)
5 X_valid = X.iloc[int(len(X)*split) : ,:].reset_index(drop=True)
6
7 y_train = y.iloc[ : int(len(X)*split)].reset_index(drop=True)
8 y_valid = y.iloc[int(len(X)*split) : ].reset_index(drop=True)
9
10 print(f"Train X size = {len(X_train)}")
11 print(f"Train y size = {len(y_train)}")
12 print(f"Valid X size = {len(X_valid)}")
13 print(f"Valid y size = {len(y_valid)}")
```

Train X size = 210
Train y size = 210
Valid X size = 90
Valid y size = 90

```
1 # Save a copy of X and y
2 # TODO might not need it
3 X_train_org = X_train.copy()
4 y_train_org = y_train.copy()
```

Polynomial Regression

a basic example

```
1 # polynomial degree
2 polynomial_degree = basic_polynomial_degree
3
```

Convert

$aX + bX^2 + cX^3 + d$

to

$aX1 + bX2 + cX3 + d$

```
1 # Add the polynomial's terms as features
2 # so that the univariate non-linear regression
3 # becomes a multivariate linear regression
4 # where every polynomial term is a feature for
5 # the linear regression
6
7 X_train_copy = X_train.copy()
8 X_valid_copy = X_valid.copy()
9 polynomial_to_linear_regression(X_train_copy, polynomial_degree)
10 polynomial_to_linear_regression(X_valid_copy, polynomial_degree)
11
12
13 X_train_copy.head()
```

	0	x	x2	x3	x4	x5	x6	x7	x8	x9	x10
0	1	-0.122797	0.015079	-0.001852	0.000227	-0.000028	3.428672e-06	-4.210306e-07	5.170128e-08	-6.348761e-09	7.796086e-10
1	1	0.954990	0.912006	0.870957	0.831756	0.794319	7.585666e-01	7.244237e-01	6.918176e-01	6.606791e-01	6.309421e-01
2	1	-0.087699	0.007691	-0.000675	0.000059	-0.000005	4.549648e-07	-3.990011e-08	3.499213e-09	-3.068787e-10	2.691306e-11
3	1	-0.405126	0.164127	-0.066492	0.026938	-0.010913	4.421212e-03	-1.791148e-03	7.256409e-04	-2.939761e-04	1.190974e-04
4	1	0.154457	0.023857	0.003685	0.000569	0.000088	1.357840e-05	2.097281e-06	3.239401e-07	5.003488e-08	7.728246e-09

Training a basic example

Normal equation for comparison :

- 0.07 training loss
- 0.09 validation loss

5k iter 10th degree polynomial lr=2.3

- Training Loss : 0.07766969752936674
- Validation Loss : 0.094857479611961

```
1 # Initialize the weights with zero
2 theta = np.array([0.0]*len(X_train_copy.columns))
3
4 # Initialize the weights with random values
5 # theta = np.random.rand(len(X_train_copy.columns),)
6
7 print("notice : takes approximately 1 minute for 1k iters (MSE)")
8
9 # tip : nice way to find decay (https://machinelearningmastery.com/using-learning-rate-schedules-deep-learning-models-python-keras/)
10 # Decay = LearningRate / Epochs
11 # Decay = 0.1 / 1000
12 # Decay = 0.0001
13 # Note : didn't work well for us :(
14
15 # Hyper-Parameters
16 iterations = basic_iterations
17 learning_rate = basic_learning_rate
18 decay = basic_decay
19
20 # Start the training
21 history, loss, theta = gradientDescent(X_train_copy,
22                                       y_train,
23                                       theta,
24                                       learning_rate,
25                                       iterations,
26                                       X_valid = X_valid_copy,
27                                       y_valid = y_valid,
28                                       loss_fn=MSE,
29                                       loss_fn_prim=MSE_prim,
30                                       decay = decay*0)
31
32
```


```
notice : takes approximately 1 minute for 1k iters (MSE)
100%                               1000/1000 [00:40<00:00, 25.08it/s]
training_loss : 0.1107 | validation_loss : 0.1162
```

Plotting the fitted polynomials

```
1 # Predicting using the learned weights(theta)
2 # Not used here but useful
3 y_hat = theta*X_valid_copy
4 y_hat = np.sum(y_hat, axis=1)
5 print(RMSE(X_valid_copy, y_valid, theta))
```


0.34090818758488145

```
1 plot_curve(X_train_copy, y_train, theta, title='Linear Regression')
2
3 print(f"Training Loss : {RMSE(X_train_copy, y_train, theta)}")
4 print(f"Validation Loss : {RMSE(X_valid_copy, y_valid, theta)}\n")
```




Plotting losses

```
1 plot_loss(history, starting_iter=20, title='Losses')
```



Plotting the learning rate

```
1 plot_lr(lr=learning_rate, iterations=iterations, decay=decay)
```



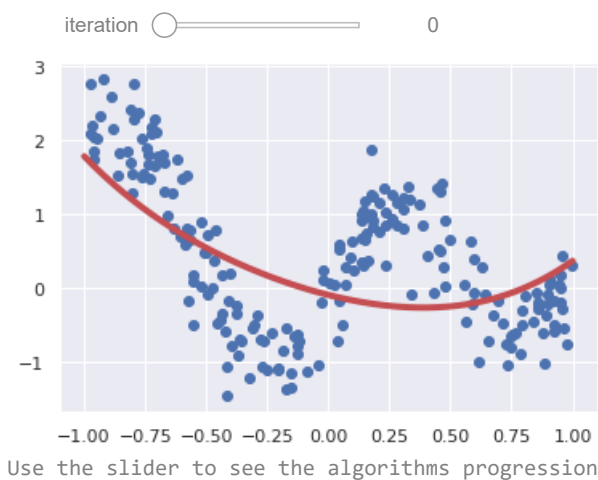
Interactive history viewer

You can use the slider to see a history of thetas

which should allow you to look at the progression

of the minization algorithm!

```
1 from ipywidgets import interact
2 import ipywidgets as widgets
3
4 @interact(iteration = widgets.IntSlider(min=0, max=iterations-1, step=10, value=0))
5 def plot_weight_history(iteration):
6     plot_curve(X_train_copy, y_train, history["weights"][iteration])
7
8 print("Use the slider to see the algorithms progression")
```



Part 4 - Plotting every Case!

Polynomial Degree :

- 5
- 8
- 10

Loss Functions :

- MSE
- RMSE
- MAE

Iterations :

- 1000
- 10000

```
1 # Global Hyper-Parameters
2 # we plot the 'max_iteration' curve (green)
3 # and 'other_iteration_to_display' curve (red)
4 iterations = max_iterations
5 other_iteration_to_display = iterations//10 # The red line in the plots ()
6
7 # Grid Hyper-Parameters
8 polynomial_degrees = [10, 8, 5]
9
10 loss_functions = [(MAE, MAE_prim),
11                  (RMSE, RMSE_prim),
12                  (MSE, MSE_prim)]
13 fn_labels = ["MAE", "RMSE", "MSE"]
14
15 #           MAE  RMSE  MSE
16 lrs       = [0.26, 1.4, 2.23]
17 decays = [0, 0, 5e-9]
```

Training all 9 models

```
1 # Containers used for storing results about each model
2 # Used later by the plotting functions
3 thetas = [[[],[],[]],
4           [[],[],[]],
5           [[],[],[]]]
6
7 losses = [[[],[],[]],
8           [[],[],[]],
9           [[],[],[]]]
10
11 histories = [[[],[],[]],
12              [[],[],[]],
13              [[],[],[]]]
14
15 # start the training process for each model
16 for i,degree in enumerate(polynomial_degrees):
17     for j, (loss_fn, loss_fn_prim) in enumerate(loss_functions):
18         print(f"degree: {degree} | loss function: {fn_labels[j]}")
19
20         # preprocess data (univariate non-linear to multivariate linear)
21         X_train_copy = X_train.copy()
22         X_valid_copy = X_valid.copy()
23         polynomial_to_linear_regression(X_train_copy, degree)
24         polynomial_to_linear_regression(X_valid_copy, degree)
25
26         # Initialize the weights with random values
27         theta = np.random.rand(len(X_train_copy.columns),)
28
29         # j(loss function) specific learning rate
30         _learning_rate = lrs[j]
31
32         # j(loss function) specific decay value
33         _decay = decays[j]
34
35         # Start the training
36         history, loss, theta = gradientDescent(X_train_copy,
37                                               y_train,
38                                               theta,
39                                               _learning_rate,
40                                               iterations,
41                                               X_valid = X_valid_copy,
42                                               y_valid = y_valid,
43                                               loss_fn = loss_fn,
44                                               loss_fn_prim = loss_fn_prim,
45                                               decay = _decay)
46
47         # Saving latest iteration's theta for each model
48         thetas[i][j].append(history["weights"][-1])
49         # Saving halfway theta for each model
50         thetas[i][j].append(history["weights"][int(other_iteration_to_display)-1]) # -1 : zero-indexed
51
52         # Saving Training loss for each model
53         losses[i][j].append(loss_fn(X_train_copy, y_train, theta))
54         # Saving Validation loss for each model
55         losses[i][j].append(loss_fn(X_valid_copy, y_valid, theta))
56         # Saving Histories for each model (used for plotting loss per iteration)
57         histories[i][j].append(history)
58
59         print()
```

degree: 10 | loss function: MAE
100% 10000/10000 [07:09<00:00, 23.95it/s]
training_loss : 0.172 | validation_loss : 0.1766

degree: 10 | loss function: RMSE
100% 10000/10000 [12:15<00:00, 13.75it/s]
training_loss : 0.2825 | validation_loss : 0.2919

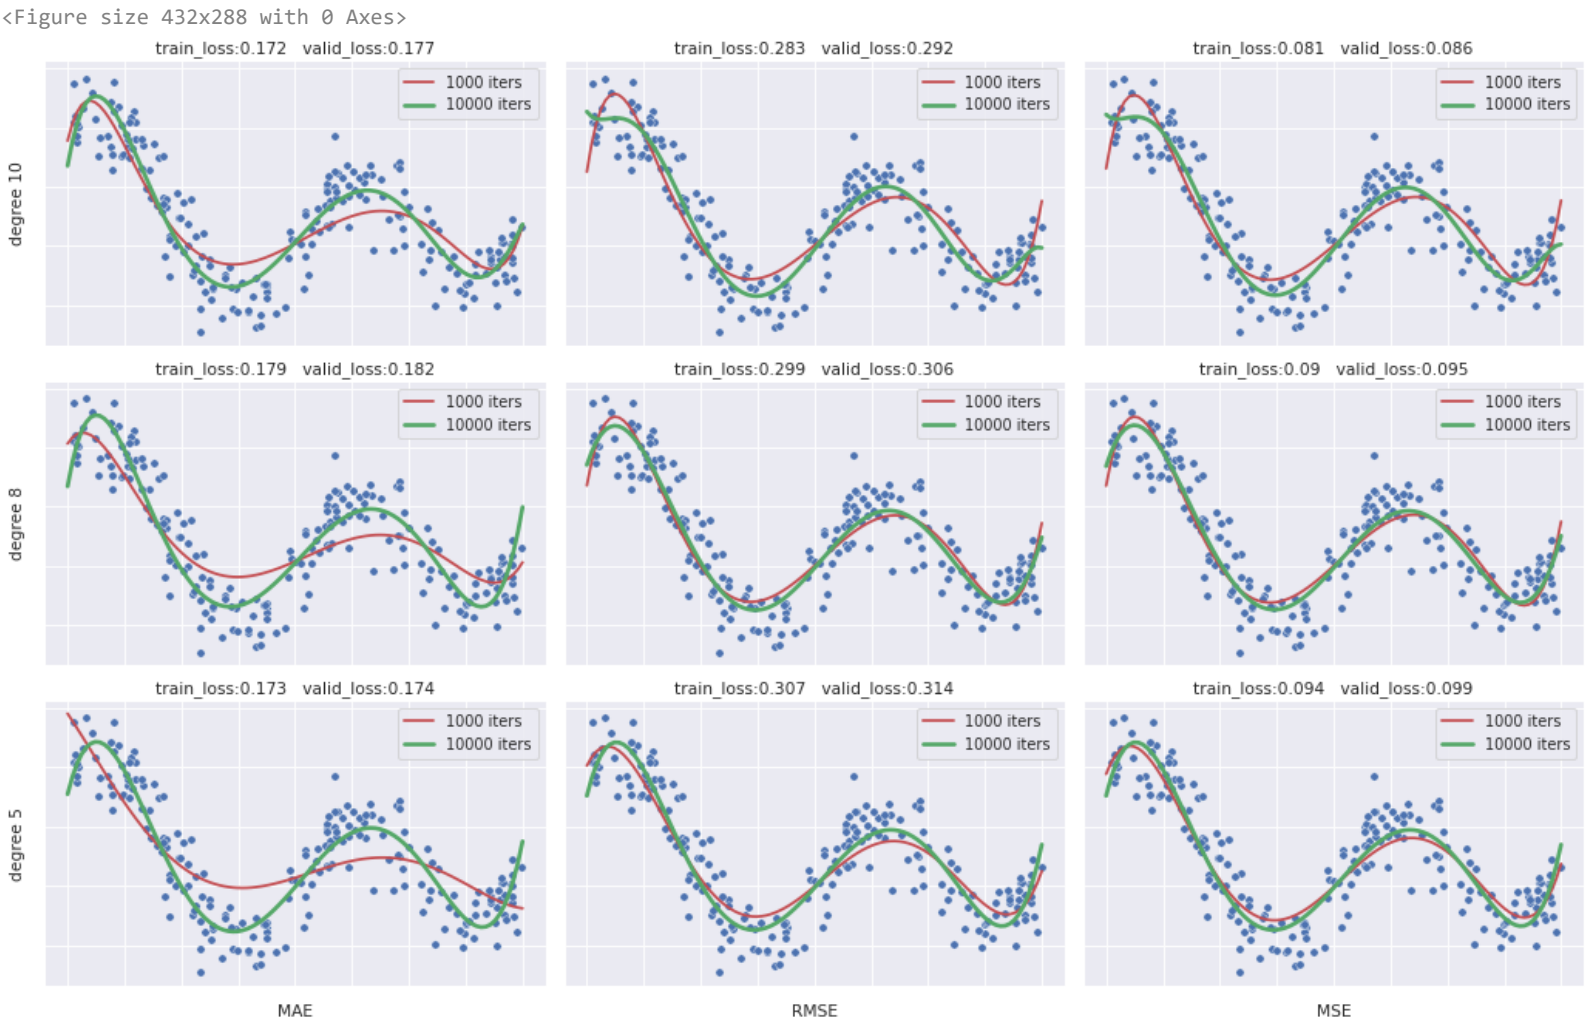
degree: 10 | loss function: MSE
100% 10000/10000 [06:42<00:00, 25.42it/s]
training_loss : 0.0807 | validation_loss : 0.0858

degree: 8 | loss function: MAE
100% 10000/10000 [05:59<00:00, 27.40it/s]
training_loss : 0.179 | validation_loss : 0.1819

Plotting fitted curves

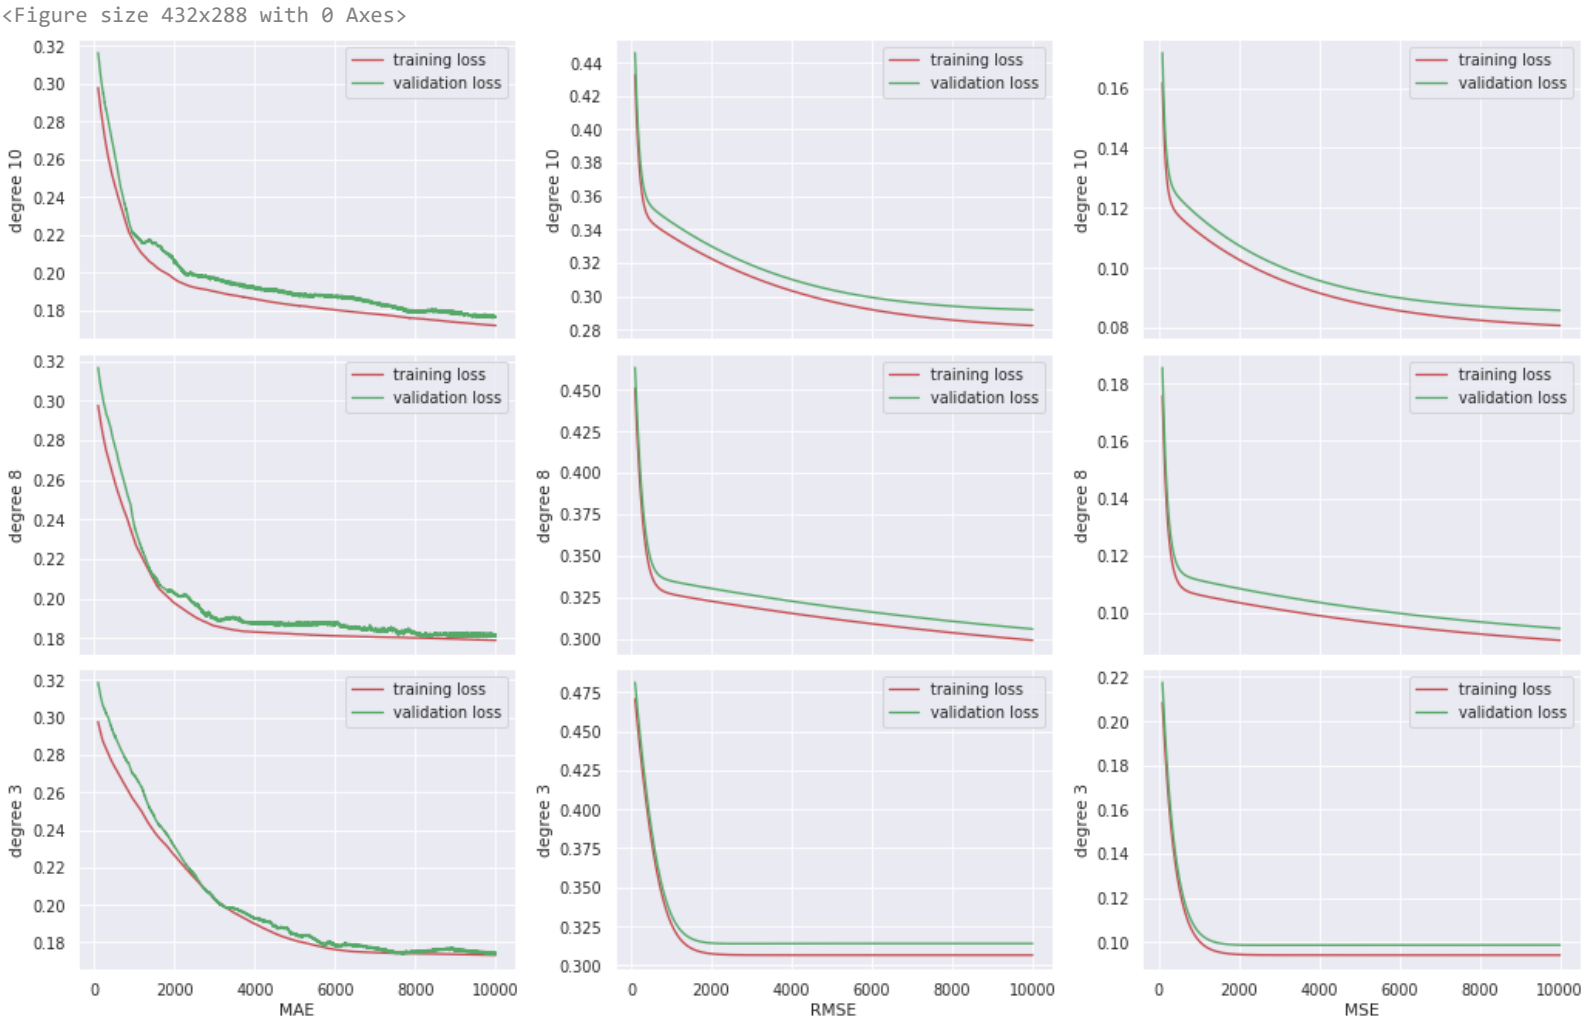
training_loss : 0.2991 | validation_loss : 0.3059

```
1 plot_every_curve(X_train,  
2 y_train,  
3 thetas,  
4 line_widths=[3,2],  
5 i_max=len(polynomial_degrees),  
6 j_max=len(loss_functions),  
7 dpi=64)
```




Plotting train/valid loss

```
1 plot_every_case_loss(histories,  
2 starting_iter=100,  
3 i_max=len(polynomial_degrees),  
4 j_max=len(loss_functions),  
5 dpi=64)
```



Plotting Learning Rate

```
1 for lr, decay, fn_name in zip(lrs, decays, fn_labels):  
2 print(f'{fn_name} | initial learning rate = {lr} | decay={decay}')  
3 plot_lr(lr=lr, iterations=iterations, decay=decay)  
4 print()
```





▼ Part 5 - Normal Equation


```
1 # preprocess data (univariate non-linear to multivariate linear)
2 X_train_copy = X_train.copy()
3 X_valid_copy = X_valid.copy()
4 polynomial_to_linear_regression(X_train_copy, 8)
5 polynomial_to_linear_regression(X_valid_copy, 8)
```

```
1 normal_theta = normalEquation(X_train_copy, y_train)
```

```
1 # Plotting the fitted polynomial over training data
2 plot_curve(X_train_copy, y_train, normal_theta, title="No regularization")
3 print(f"Training Loss(RMSE) : {RMSE(X_train_copy, y_train, normal_theta)}")
```



```
1 # Plotting the fitted polynomial over validation data
2 plot_curve(X_valid_copy, y_valid, normal_theta, title="No regularization")
3 print(f"Validation Loss : {RMSE(X_valid_copy, y_valid, normal_theta)}\n")
```



▼ Part 6 - Regularized Normal Equation

```
1 loss_fn = RMSE # as asked in the question
2
3 reg_normal_theta1 = regularizedNormalEquation(X_train_copy, y_train, lambd=0.075)
4 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=0.1")
5 reg_normal_theta2 = regularizedNormalEquation(X_train_copy, y_train, lambd=0.75)
6 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=1")
7 reg_normal_theta3 = regularizedNormalEquation(X_train_copy, y_train, lambd=7.5)
8 # plot_curve(X_train, y_train, reg_normal_theta, title="lambda=4")
```

```
1 names = ["λ = 0.075", "λ = 0.75", "λ = 7.5"]
2
3 training_errors = [round(loss_fn(X_train_copy, y_train, reg_normal_theta1),5),
4                    round(loss_fn(X_train_copy, y_train, reg_normal_theta2),5),
5                    round(loss_fn(X_train_copy, y_train, reg_normal_theta3),5)]
6
7 validation_errors =[round(loss_fn(X_valid_copy, y_valid, reg_normal_theta1),5),
8                    round(loss_fn(X_valid_copy, y_valid, reg_normal_theta2),5),
9                    round(loss_fn(X_valid_copy, y_valid, reg_normal_theta3),5)]
10
11 for i in range(len(names)):
12     print(f"{i} - Regularized Normal Equation ({names[i]})")
13     print(f"Training Loss : {training_errors[i]}")
14     print(f"Validation Loss : {validation_errors[i]}\n")
```

0 - Regularized Normal Equation (λ = 0.075)
Training Loss : 0.33571
Validation Loss : 0.34334

1 - Regularized Normal Equation (λ = 0.75)
Training Loss : 0.43287
Validation Loss : 0.44485

2 - Regularized Normal Equation (λ = 7.5)
Training Loss : 0.51534
Validation Loss : 0.53339

```
1 plot_normal_equations(X_train, y_train, dpi=105)
```

Thanks for reading. if you have any questions:

- gholamrezadar@gmail.com
- instagram : @gholamreza_dar
- telegram : @gholamrezadar

